

# Kubernetes for DevOps



Nguyen Canh Hat

June 2023

**Nash  
Tech.**

# Agenda

- 1. Kubernetes Overview**
  - Benefits of Kubernetes
  - What is Kubernetes?
- 2. Kubernetes Architecture**
- 3. Primary Kubernetes Objects**
  - Deployment, StatefulSet, DaemonSet
  - Probes, Deployment Strategies
  - Resource Request and Limit
  - CronJob
  - Network
  - Configuration data
  - Persistent Storage
- 4. Deployment Strategies**
- 5. Kubernetes User Interaction**
  - K8s CLI
  - External tool

# Kubernetes Overview



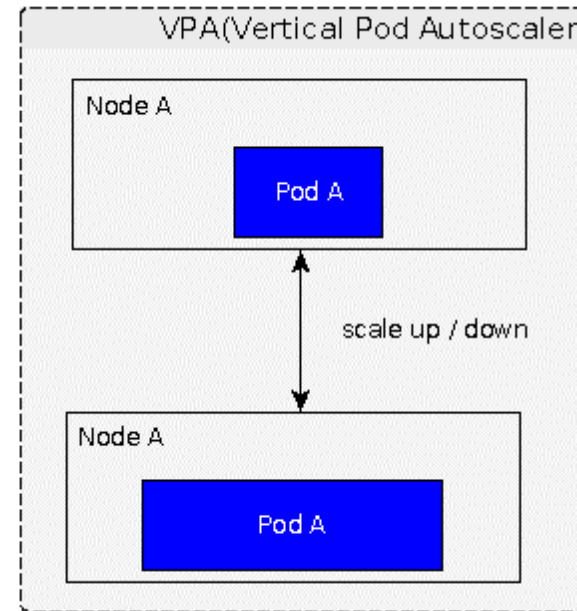
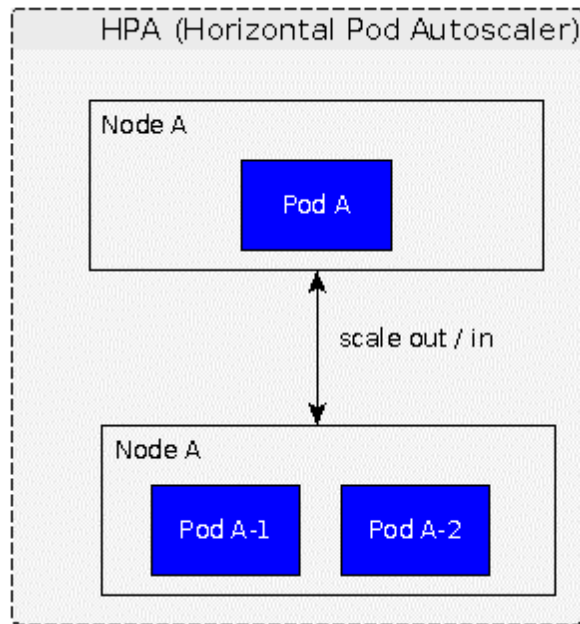
# Benefits of Kubernetes – High Availability (1/3)

## Docker

- NOT able to manage auto-scaling itself

## Kubernetes – auto scaling

- Provide HPA, VPA, CA



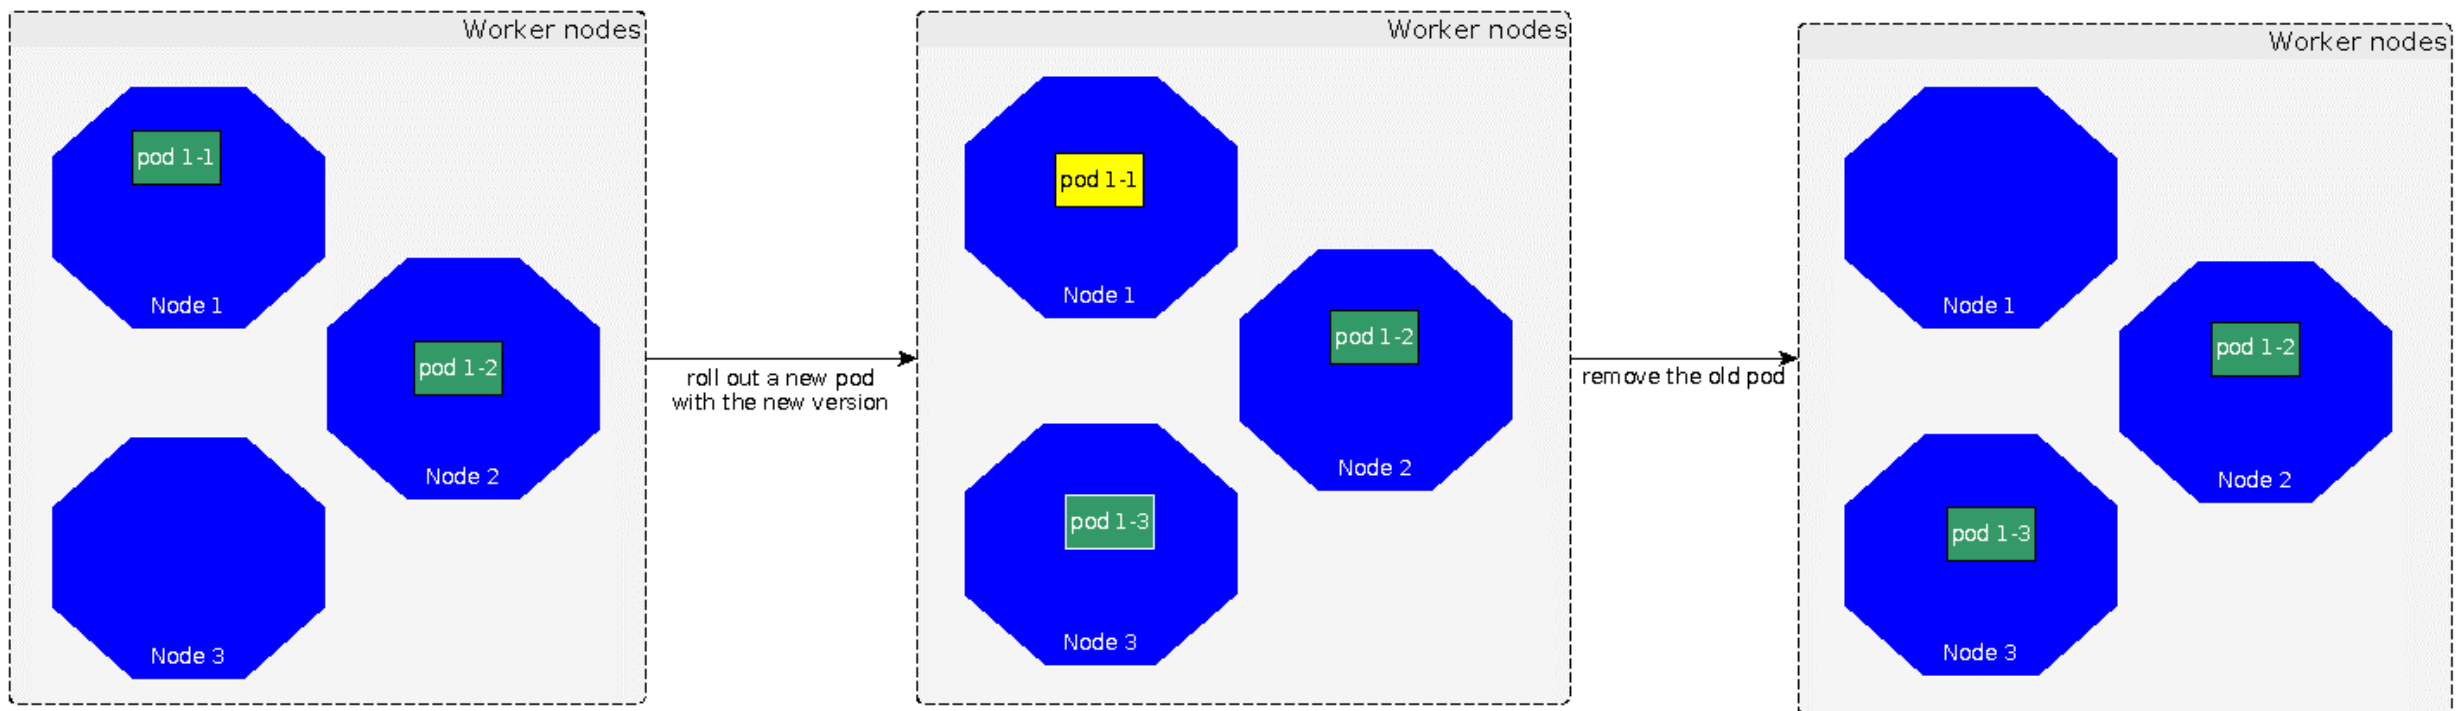
# Benefits of Kubernetes – High Availability (2/3)

## Docker

- Introduce a downtime where there is an update

## Kubernetes – zero downtime deployment

- Deployment strategies: rolling update, recreate



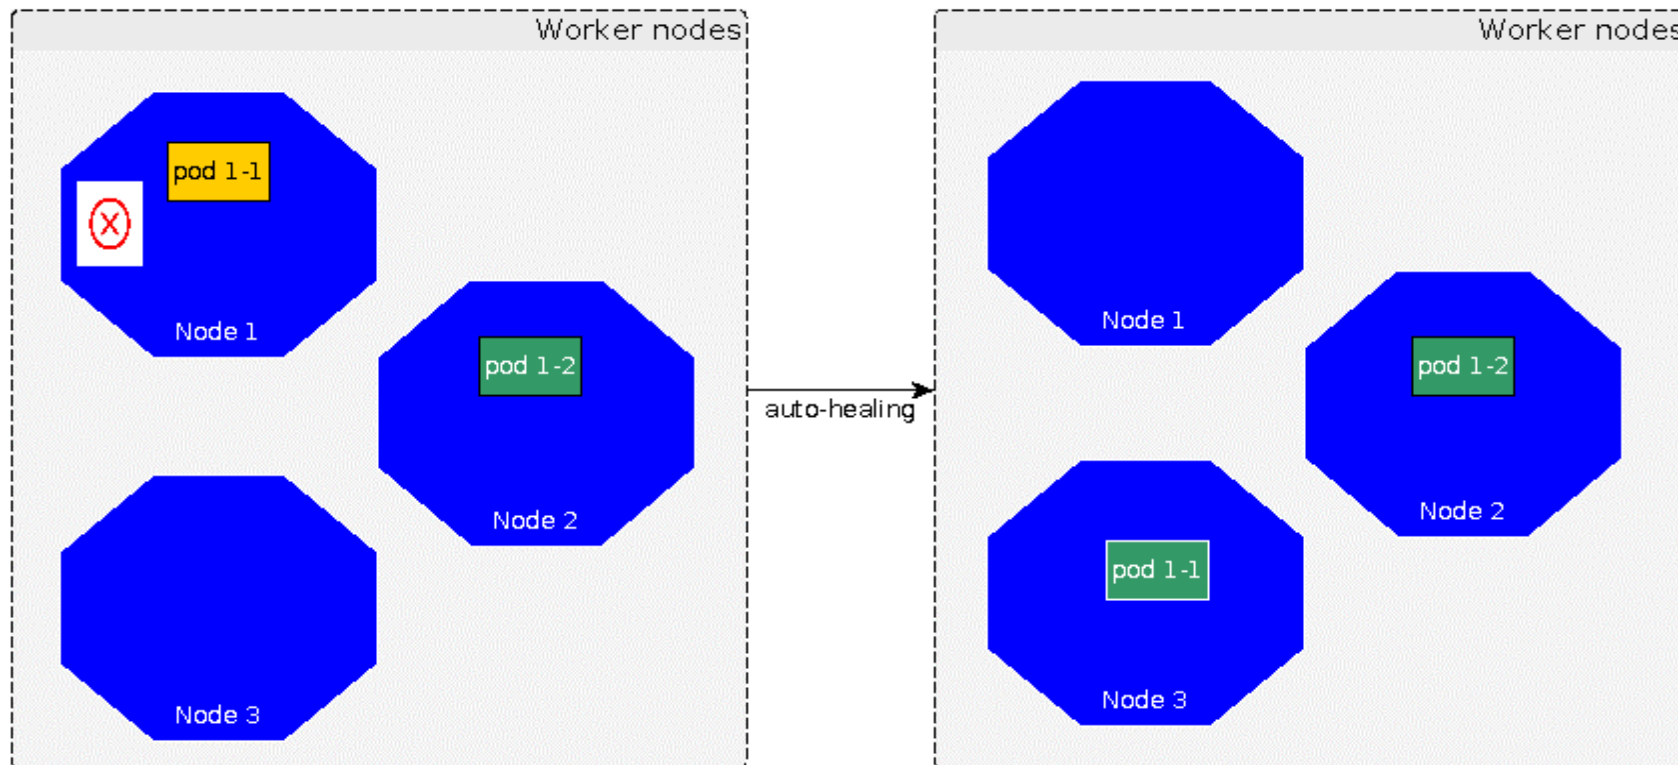
# Benefits of Kubernetes – High Availability (3/3)

## Docker

- NOT able to get failed applications back

## Kubernetes – auto healing capability

- Probes: detect failure and recover applications



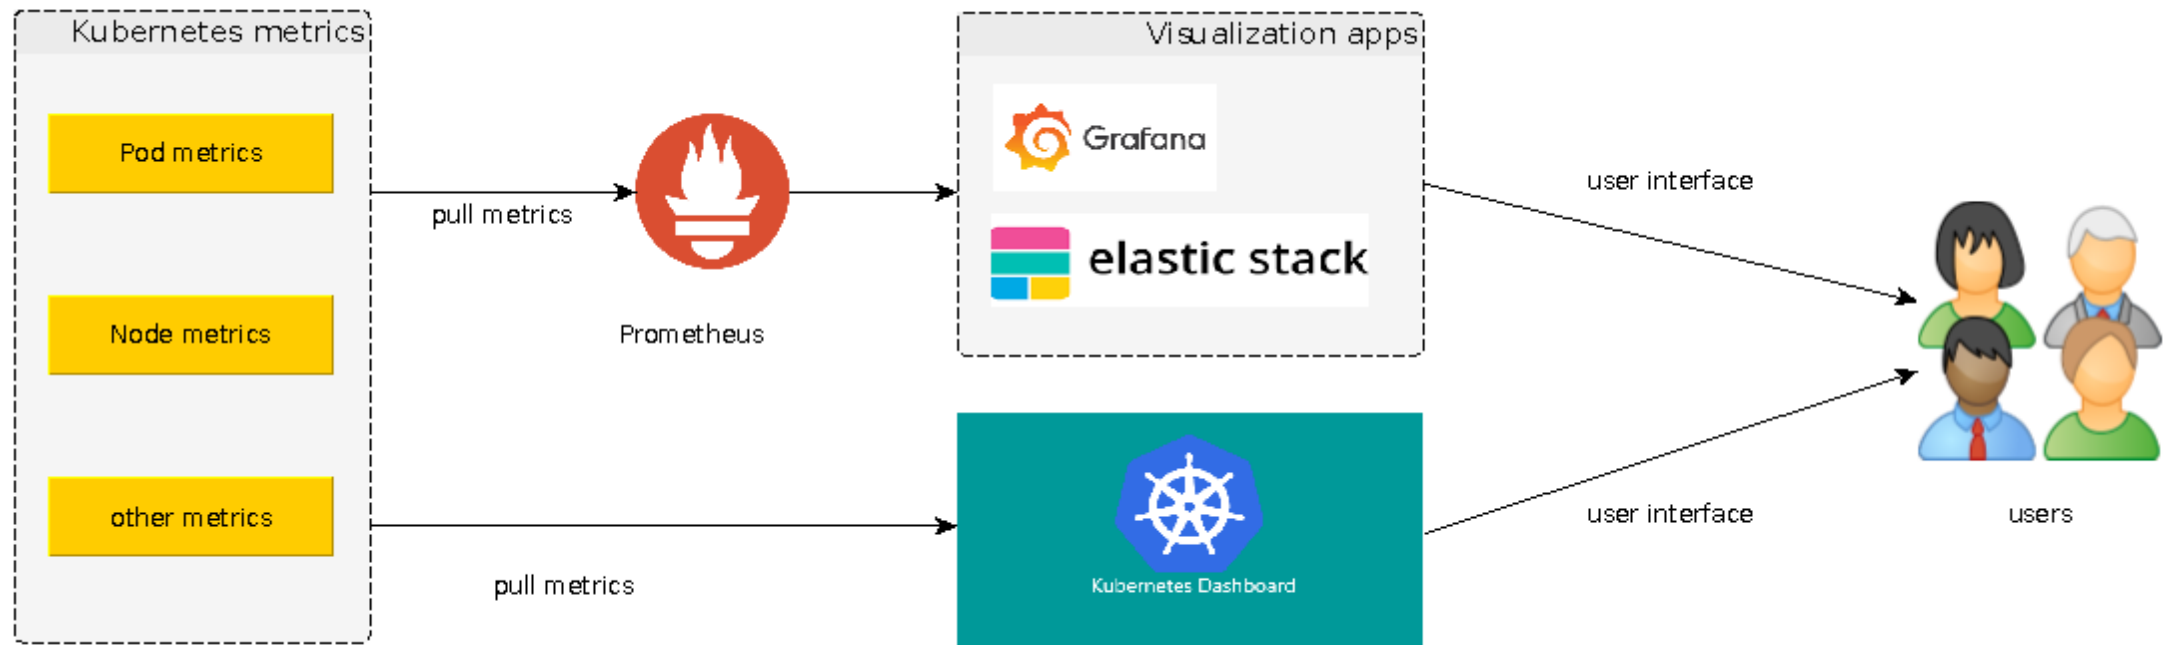
# Benefits of Kubernetes – Monitoring

## Docker

- Challenges of setting up monitoring tools

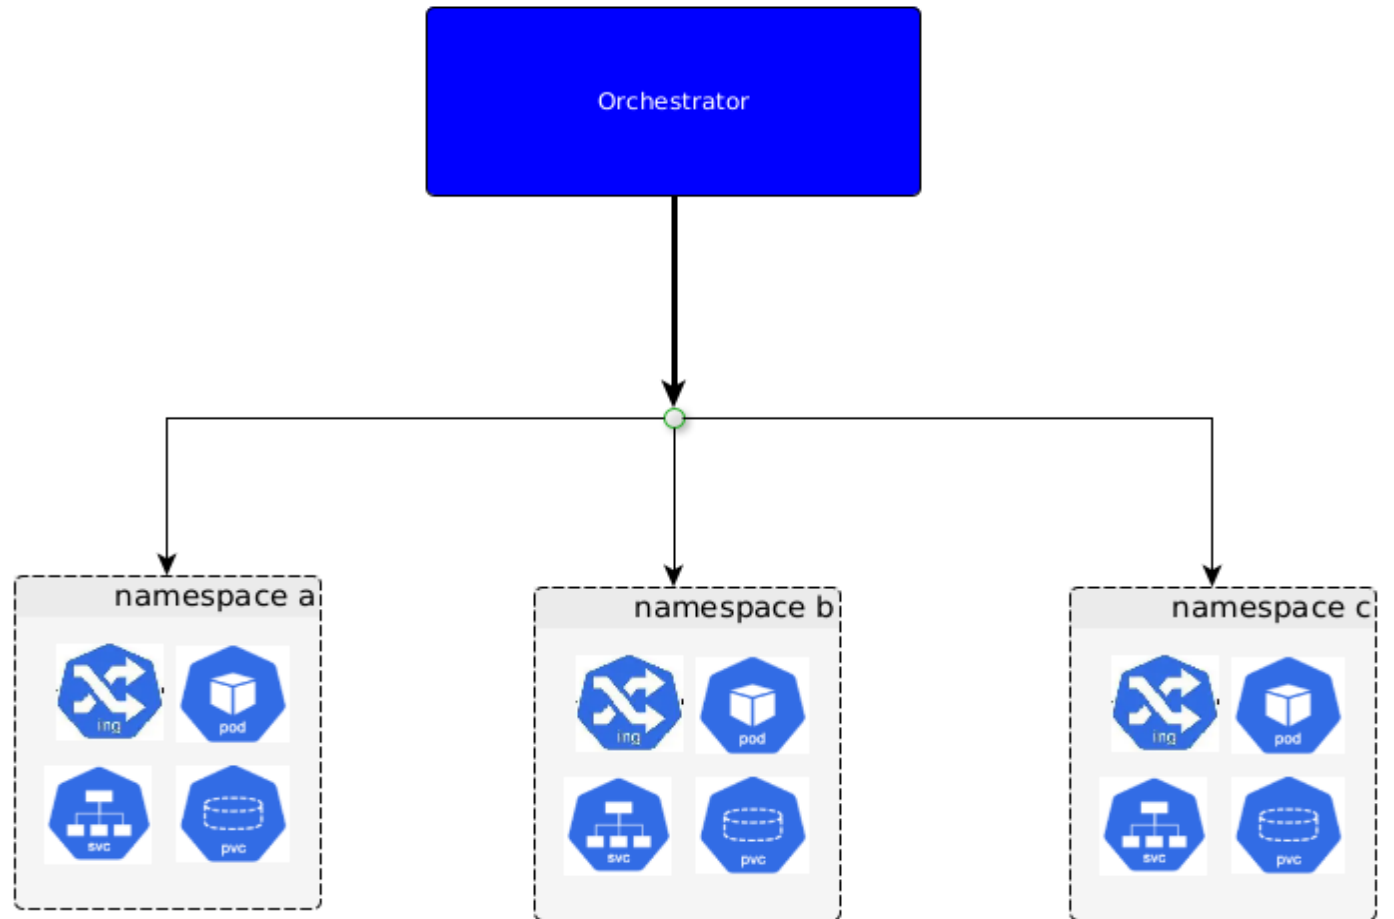
## Kubernetes – Kubernetes APIs

- Provide metrics and built in functions to integrate with 3<sup>rd</sup> tools



# What is Kubernetes?

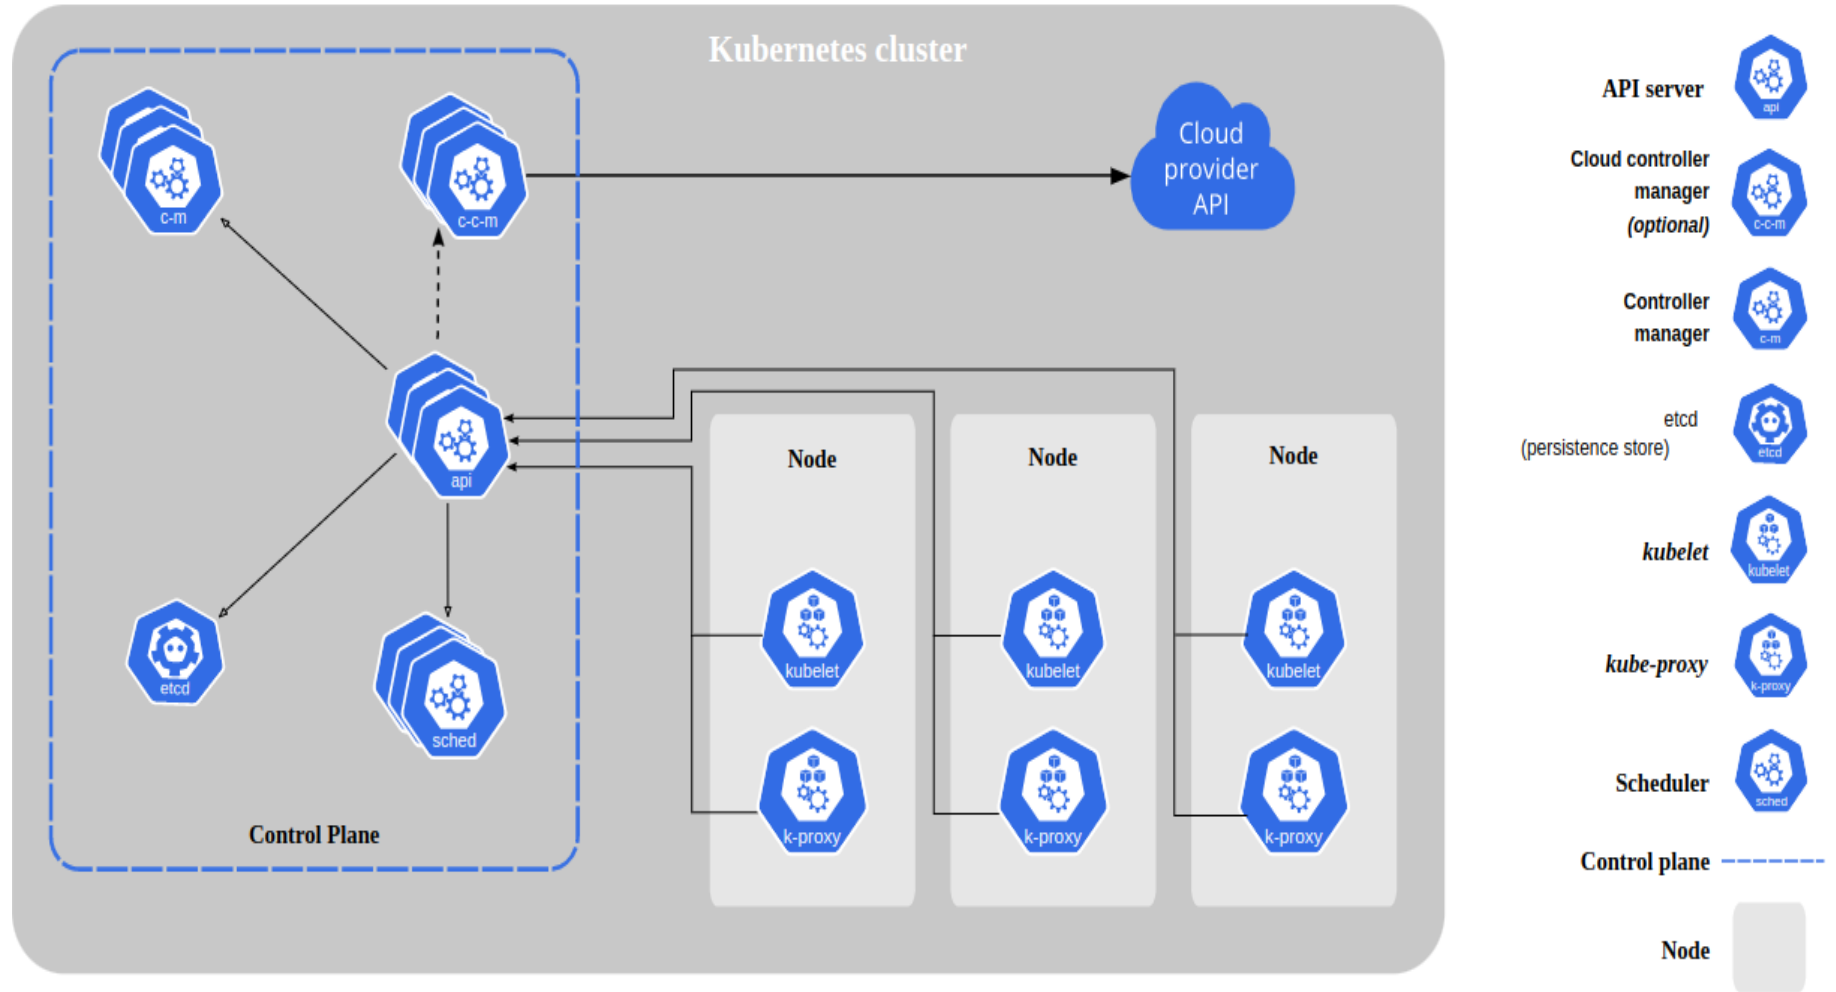
- Found by Google
- Open source
- Container orchestrator
- Automate manual processes (deployment, scaling, management of containerized applications)





# Kubernetes Architecture

- Control plane (master)
- Nodes (worker)

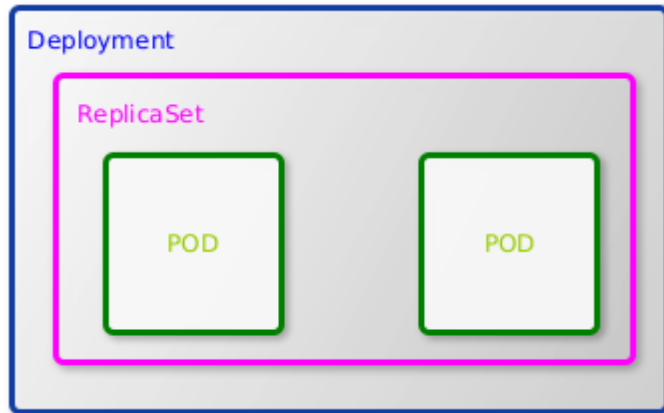


# Primary Kubernetes Objects



# Deployment

- A declarative template for pods and replicaSet
- Used for **stateless** application



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

# StatefulSet

Similar to “deployment” but:

- Used for **stateful** application
- There are 3 sections that can be updated on the fly:
  - replicas
  - template
  - updateStrategy
- Ordered pod creation/deletion
- Pod can be accessed by its own service

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: postgres-sts
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgresql
          image: postgresql:latest
          imagePullPolicy: "IfNotPresent"
          env:
            - name: POSTGRES_USER
              value: "postgres"
            - name: POSTGRES_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: postgres-secrets
                  key: postgresql-password
            - name: POSTGRES_DB
              value: "postgres"
          ports:
            - name: postgresql
              containerPort: 5432
              protocol: TCP
          volumeMounts:
            - name: data-vct
              mountPath: /var/lib/postgresql/data
```

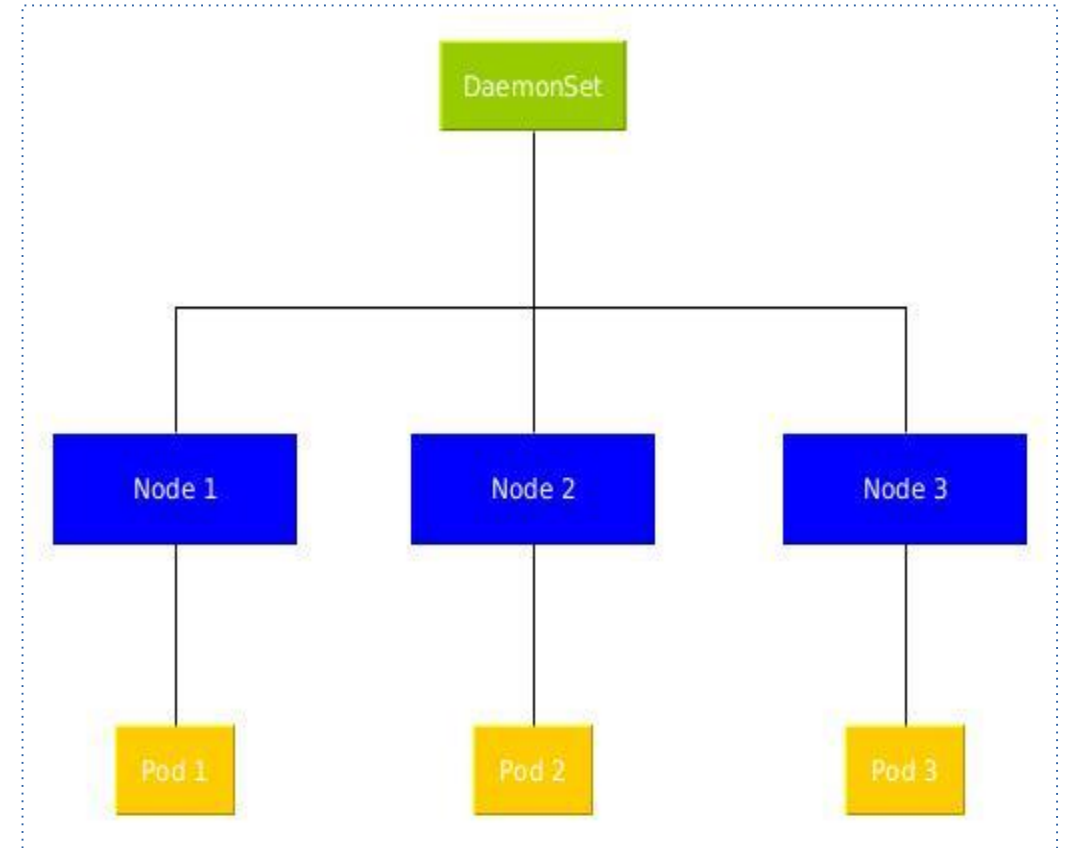
# DaemonSet

Similar to “deployment” but:

- Every node has one pod (master & worker)
- Each Pod lifetime depends on its node

Use cases:

- Install a monitoring agent per node
- Install a log connector on each node



# Probes

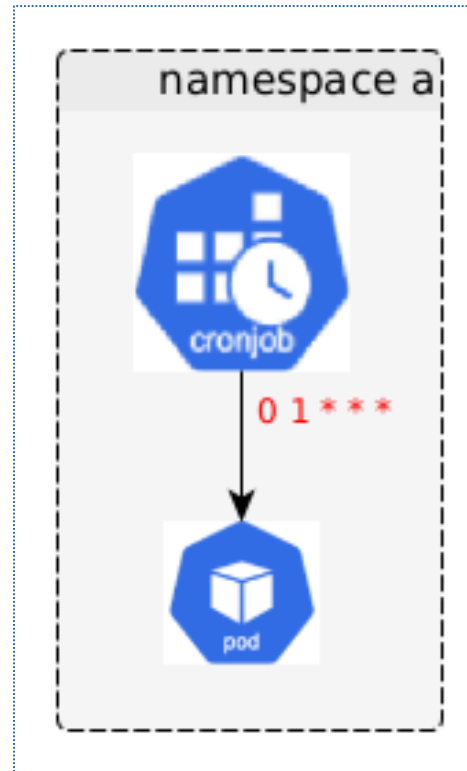
Probes are health checks that can monitor a container's status and act on it.

- StartupProbes: Indicate when the pod is **started**.
- ReadinessProbes: indicate when to start **serving**
- LivenessProbes: indicate when to **restart** the pod

```
spec:
  containers:
  - name: nginx
    image: nginx:1.14.2
    ports:
    - containerPort: 80
    livenessProbe:
      httpGet:
        path: /
        port: 80
        httpHeaders:
      initialDelaySeconds: 5
      timeoutSeconds: 1
      periodSeconds: 10
      successThreshold: 1
      failureThreshold: 3
```

# Cronjob

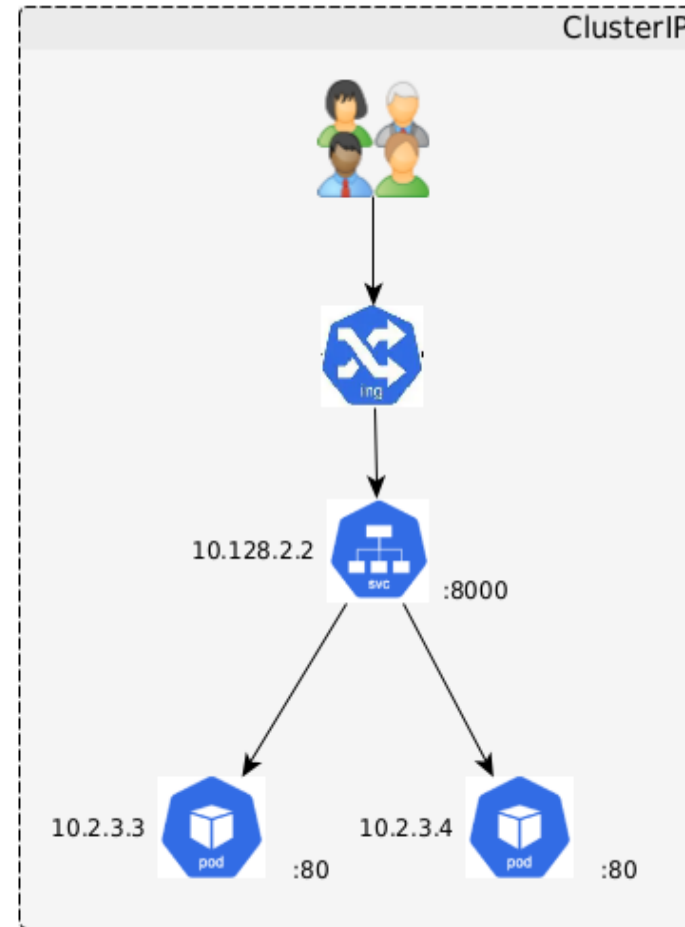
- Automatic action triggered in every scheduled time



```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: hello-friends
spec:
  schedule: "0 1 * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox:1.28
              imagePullPolicy: IfNotPresent
              command:
                - /bin/sh
                - -c
                - date; echo Hello friends
          restartPolicy: OnFailure
```

# Network – Service (1/2)

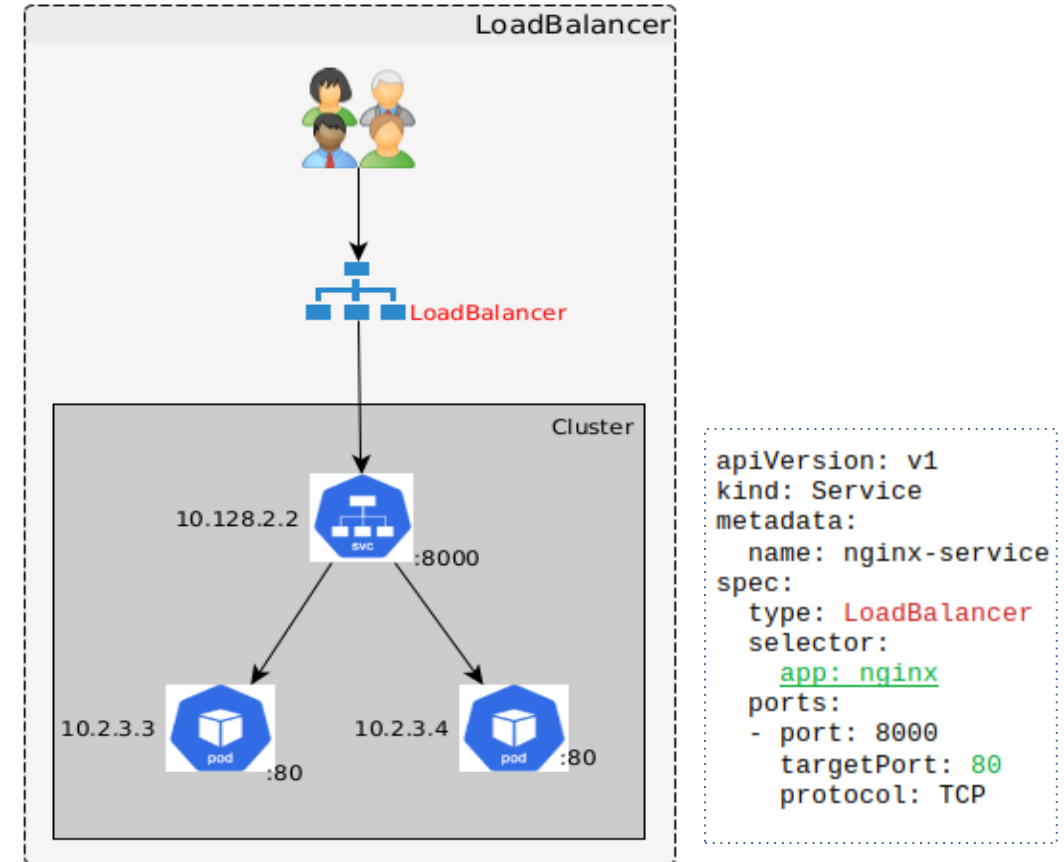
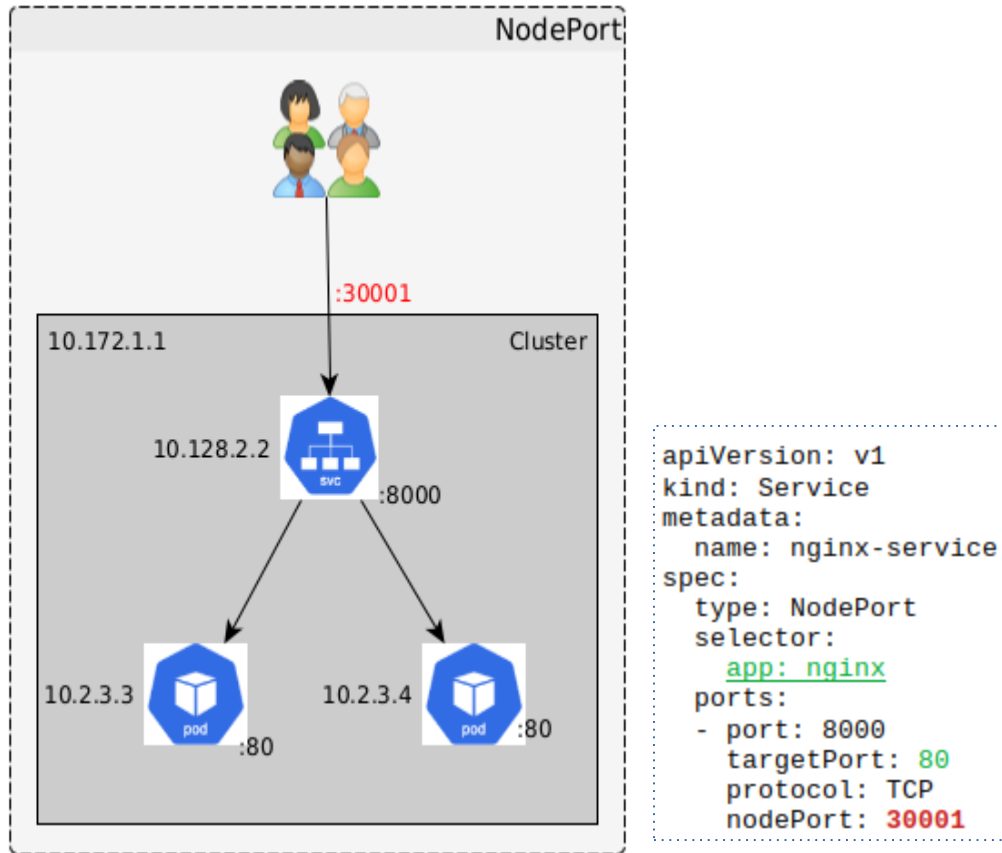
- A logical collection of pods
- Roles:
  - Enable network access to pods
  - Load balancing
- Types:
  - ClusterIP: provide internal access inside cluster only
  - NodePort:
    - Provide a unique fixed cluster node port for external access
    - Not recommended due to security (expose access to cluster)
  - LoadBalancer: Integrate with cloud providers for external access



```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - port: 8000
      targetPort: 80
      protocol: TCP
```

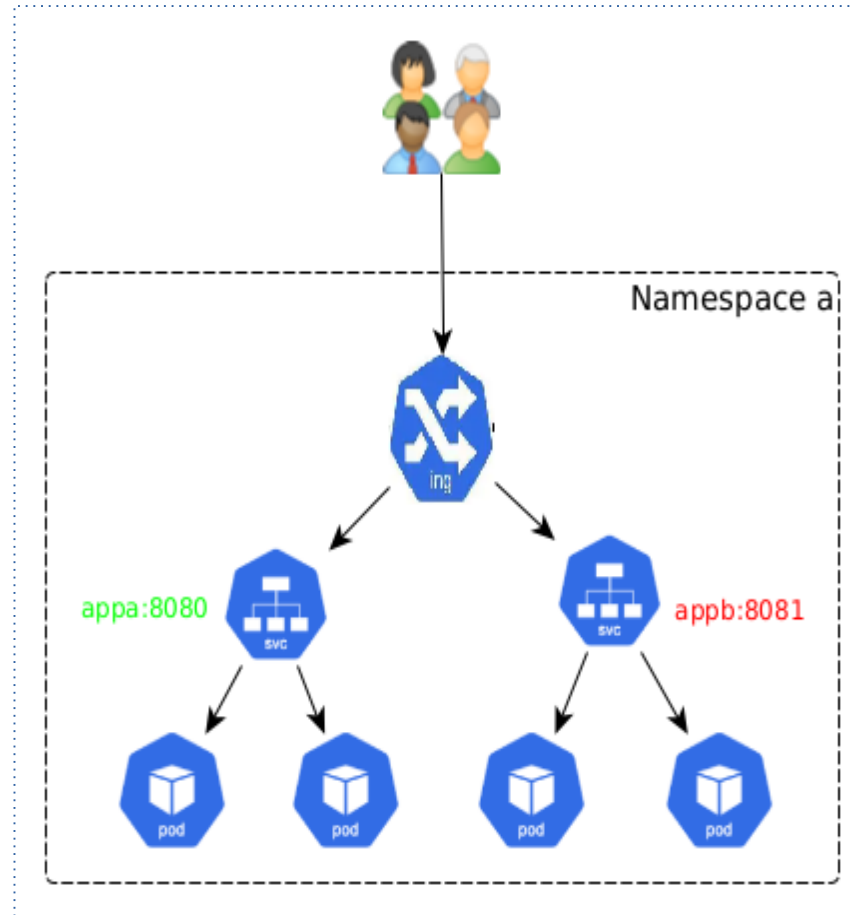


# Network – Service (2/2)



# Network - Ingress

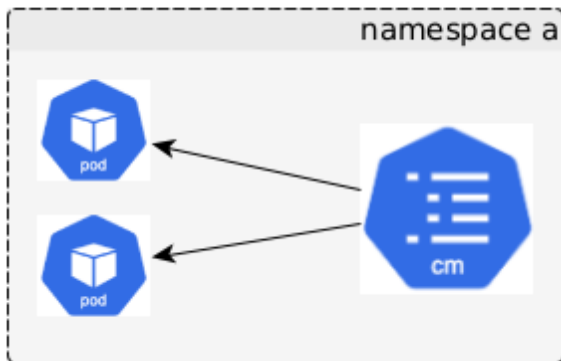
- Provide access to application by providing **readable** human URL with http/https protocols and routing **rules**.



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: app-ingress
  annotations:
    nginx.ingress.kubernetes.io/ssl-redirect: "false"
    nginx.ingress.kubernetes.io/force-ssl-redirect: "false"
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  tls:
  - hosts:
    - abc.com
    secretName: tls-secret
  rules:
  - host: abc.com
    http:
      paths:
      - path: /appa
        backend:
          serviceName: appa-service
          servicePort: 8080
      - path: /appb
        backend:
          serviceName: appb-service
          servicePort: 8081
```

# Configuration data - ConfigMap

- Key-value pairs of data
- Manage pods configuration data
  - Environment variables
  - Configuration files
- Used for **non-sensitive** data (plain text)



```
apiVersion: v1
kind: ConfigMap
metadata:
  name: nginx-config
data:
  proxy-connect-timeout: "10s"
  proxy-read-timeout: "10s"
  client-max-body-size: "2m"
```

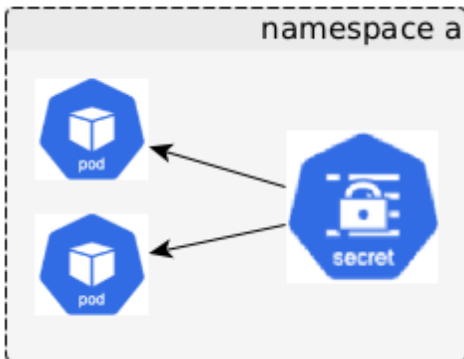
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
labels:
  app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          envFrom:
            - configMapRef:
                name: nginx-config
```

# Configuration data - Secret

- Key-value pairs of data
- Manage pods configuration data
  - Environment variables
  - Secret files
- Used for **sensitive** data (base64-encoded)

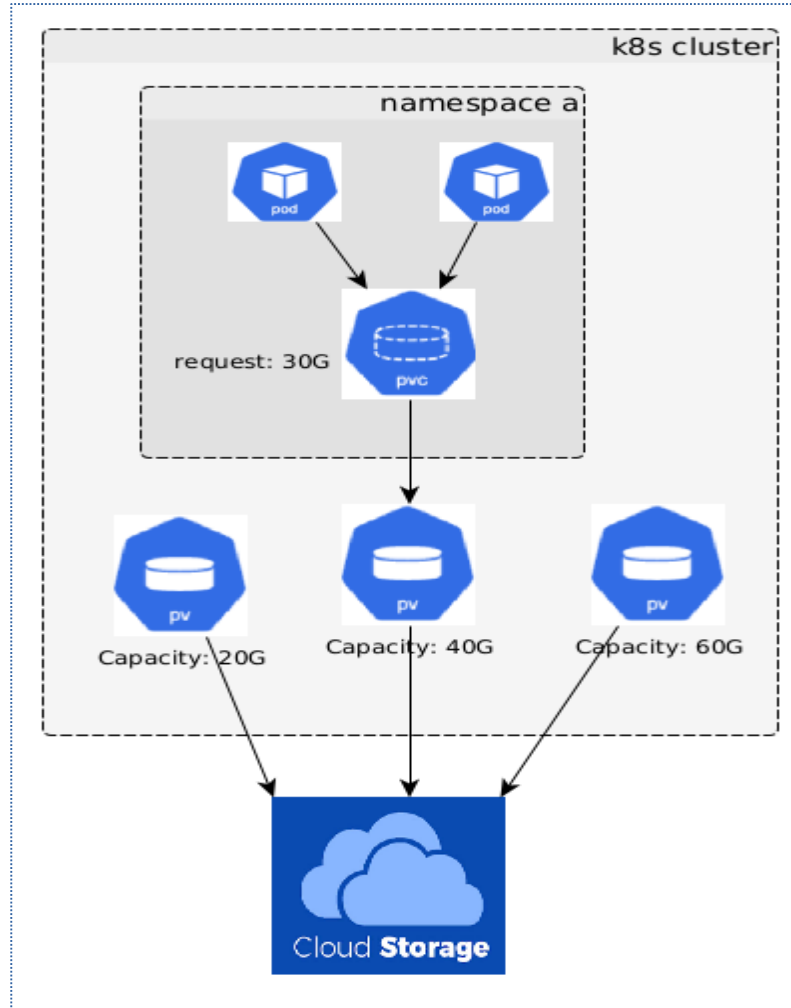
```
apiVersion: v1
kind: Secret
metadata:
  name: nginx-secrets
type: Opaque
data:
  PASSWORD: cGFzc3dvcmQK
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
labels:
  app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          envFrom:
            - configMapRef:
                name: nginx-config
              - secretRef:
                  name: nginx-secrets
```



# Persistent Storage – Persistent Volume, Persistent Volume Claim

- Persistent volume:
  - **Static** storage allocated in a cluster
  - Independent from pod's lifecycle
  - Provisioned by cluster admin
- Persistent volume claim:
  - Dynamic storage request
  - Depending on pod's lifecycle
  - Asked by pods



```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: efs-pv
spec:
  capacity:
    storage: 5Gi
  volumeMode: Filesystem
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  storageClassName: ""
  csi:
    driver: efs.csi.aws.com
    volumeHandle: [FileSystemId]::[AccessPointId]

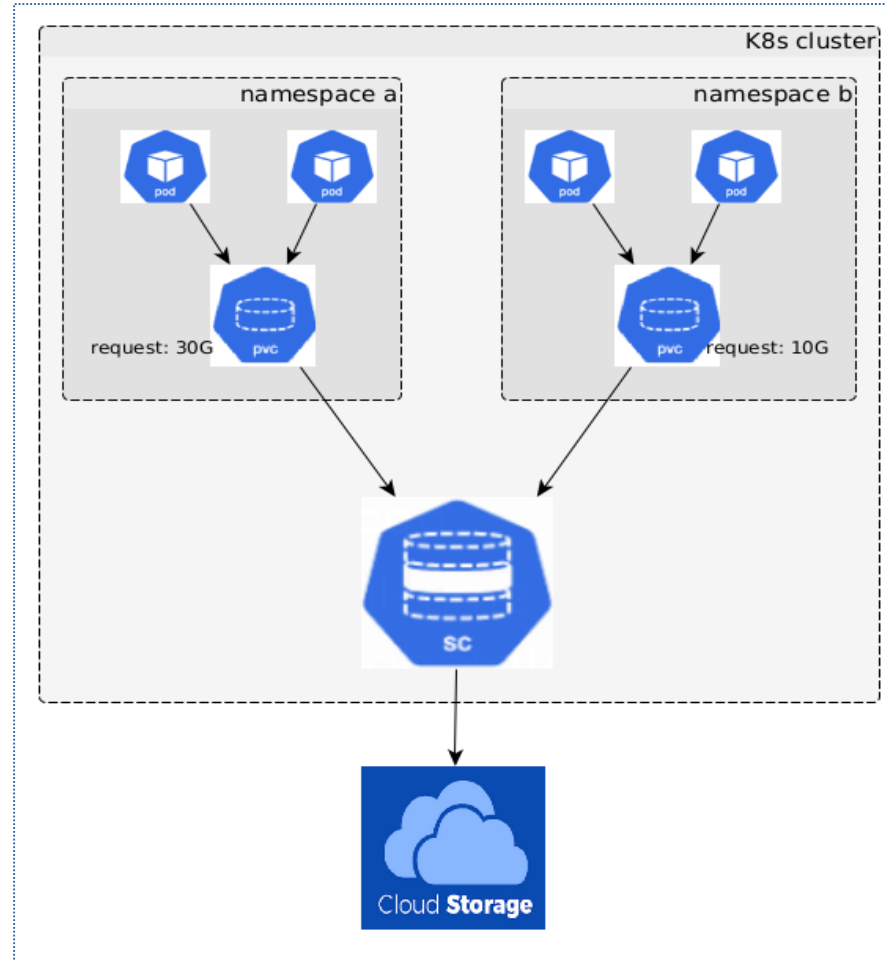
---
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: efs-claim
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: ""
  resources:
    requests:
      storage: 5Gi
```

# Persistent Storage – Persistent Volume, Persistent Volume Claim

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          envFrom:
            - configMapRef:
                name: nginx-config
            - secretRef:
                name: nginx-secrets
          volumeMounts:
            - mountPath: /var/www/html
              name: data-volume
          ports:
            - containerPort: 80
      volumes:
        - name: data-volume
          persistentVolumeClaim:
            claimName: efs-claim
```

# Persistent Storage – Storage Class

- Similar to PV, but providing a way to **dynamically** provision PV in a cluster



```
apiVersion: storage.k8s.io/v1
kind: Storage-class
metadata:
  name: ebs-sc
provisioner: kubernetes.io/aws-ebs
parameters:
  type: gp2
volumeBindingMode: WaitForFirstConsumer
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: ebs-pvc
spec:
  storageClassName: ebs-sc
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 5Gi
```

# Persistent Storage – Volume Claim Template

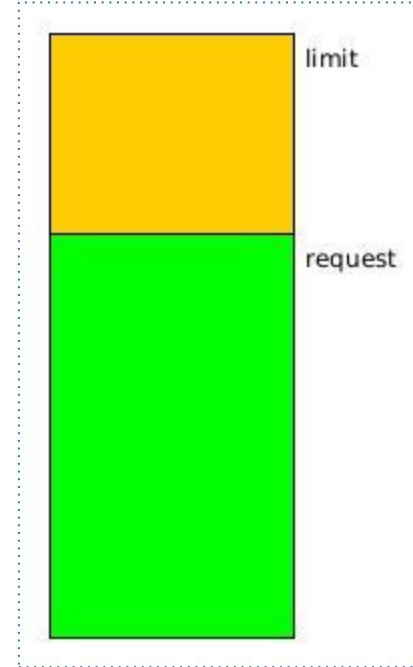
- Similar to PVC but:
- Used for stateful app (DB)
- Each pod can have its own allocated PVC

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: postgres-sts
spec:
  replicas: 3
  template:
    metadata:
      labels:
        app: postgres
    spec:
      containers:
        - name: postgresql
          image: postgresql:latest
          imagePullPolicy: "IfNotPresent"
          env:
            - name: POSTGRES_USER
              value: "postgres"
            - name: POSTGRES_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: postgres-secrets
                  key: postgresql-password
            - name: POSTGRES_DB
              value: "postgres"
          ports:
            - name: postgresql
              containerPort: 5432
              protocol: TCP
          volumeMounts:
            - name: data-vct
              mountPath: /var/lib/postgresql/data
      volumeClaimTemplates:
        - metadata:
            name: data-vct
          spec:
            storageClassName: ebs-sc
            accessModes:
              - "ReadWriteOnce"
            resources:
              requests:
                storage: "2Gi"
```



# Resource Request & Limit

- There are 2 main types of resources: CPU & memory
- Request: Ensure the minimum number of resources required for the container
- Limit: Is the maximum number of resources that Kubernetes will allow the container to use



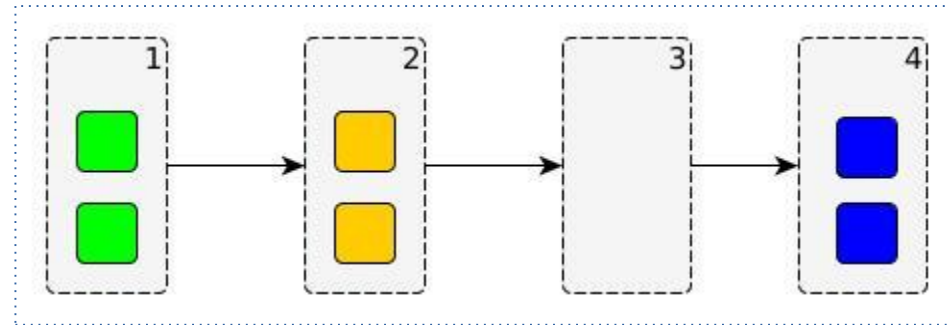
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          resources:
            limits:
              memory: 300Mi
              cpu: 1
            requests:
              memory: 100Mi
              cpu: 100m
```

# Deployment Strategies



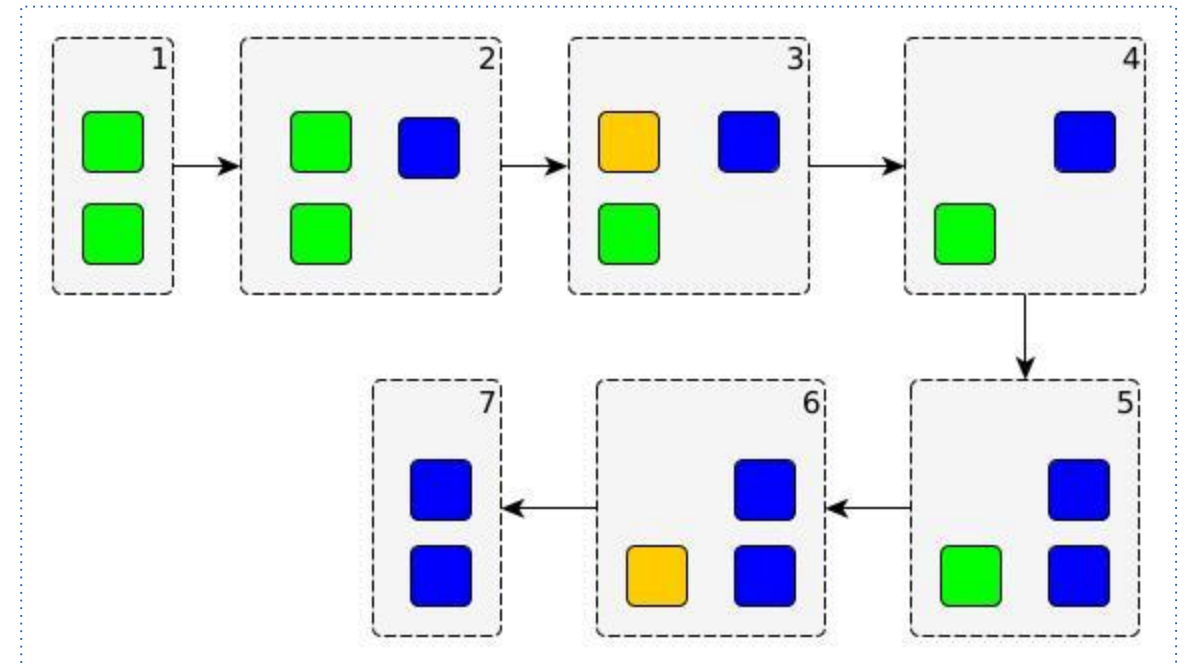
# Recreate

- Terminate all current pods, create new ones
- Use case:
  - Not able to have 2 different versions of app installed at the same time (the big update)
- Cons:
  - Introduce a downtime period.
- Pros:
  - Quick update



# RollingUpdate

- Create a few of new pods, then terminate old ones correspondingly
- Use case:
  - There is no conflict between old and new versions
- Cons:
  - There is a period when the old and new application versions serve users at the same time
- Pros:
  - No downtime



# Kubernetes User Interaction



# Command line - kubectl

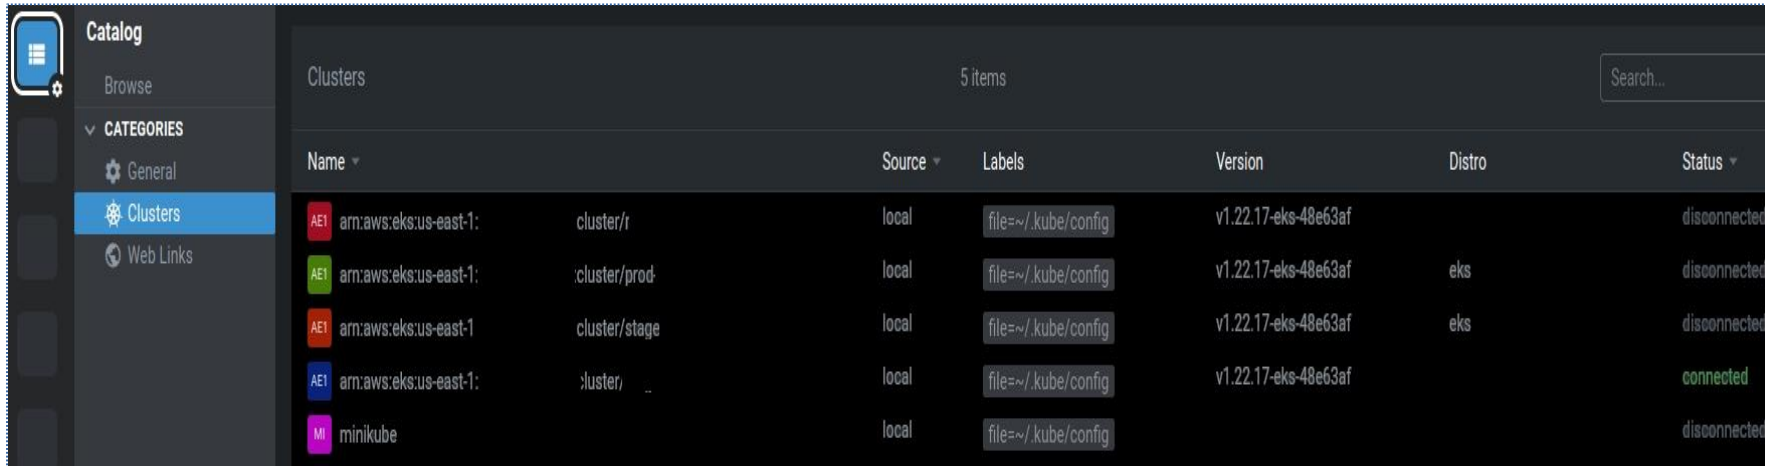
- Communicate with K8s resource
- Command line tool (refer to “[install kubectl](#)” for the installation”

```
Kubectl [action] [resource] [options]
kubectl get pod -o wide
kubectl get deployment -o wide

kubectl [action] [object] [object-name]
kubectl describe pod [pod-name]
kubectl describe deployment [deployment-name]
kubectl delete deployment [deployment-name]
```

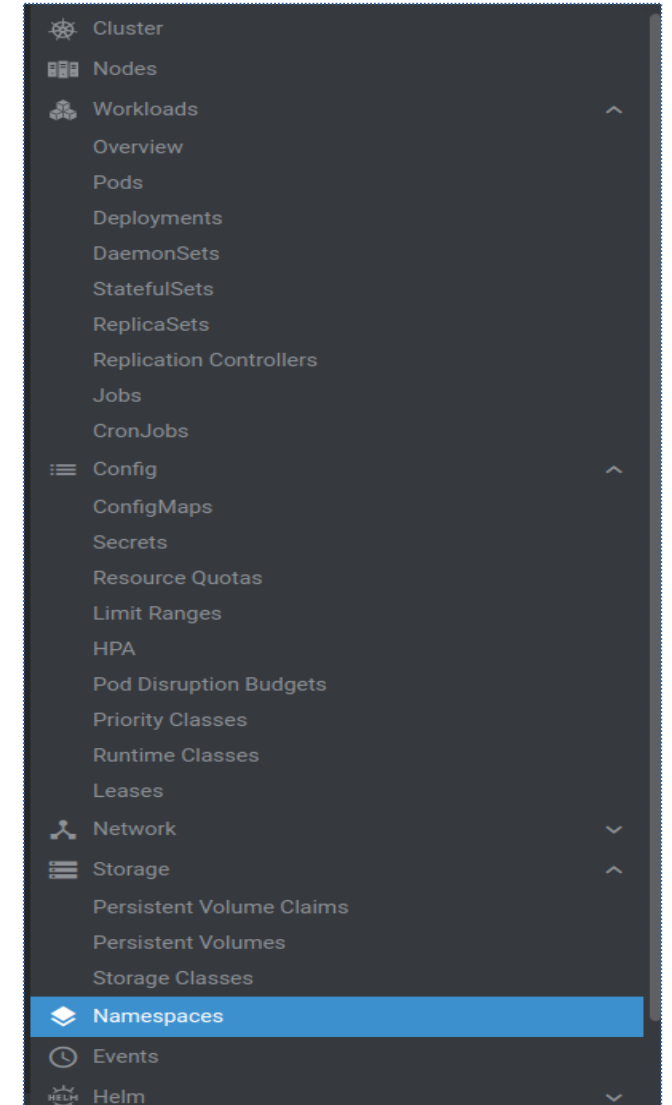
# User interface - Lens

- Communicate with K8s resource
- User interface tool (refer to “[install lens](#)” for the installation)



The screenshot shows the 'Catalog' page in the Lens UI. On the left, there is a sidebar with 'Catalog' at the top, followed by 'Browse', 'CATEGORIES', 'General', 'Clusters' (which is highlighted), and 'Web Links'. The main area displays a table of clusters. The table has columns for Name, Source, Labels, Version, Distro, and Status. There are 5 items in the list. The first four items are AWS EKS clusters, and the last one is minikube.

Name	Source	Labels	Version	Distro	Status
arn:aws:eks:us-east-1:cluster/t	local	file=~/.kube/config	v1.22.17-eks-48e63af		disconnected
arn:aws:eks:us-east-1:cluster/prod	local	file=~/.kube/config	v1.22.17-eks-48e63af	eks	disconnected
arn:aws:eks:us-east-1:cluster/stage	local	file=~/.kube/config	v1.22.17-eks-48e63af	eks	disconnected
arn:aws:eks:us-east-1:cluster/	local	file=~/.kube/config	v1.22.17-eks-48e63af		connected
minikube	local	file=~/.kube/config			disconnected



**Thank you**