Kafka for Java Developers

Sep-2022

Nash Tech.

Agenda

Understand:

- What is event streaming?
- What is Apache Kafka?
- Why Kafka?
- Delivery guarantees & Throughput

Kafka:

- Architecture and Components

Advanced:

- Kafka Connect
- Kafka Stream
- Security
- Kafka Pros-Cons

What is event streaming?

What is event streaming?

Event streaming is the digital equivalent of the human body's central nervous system.

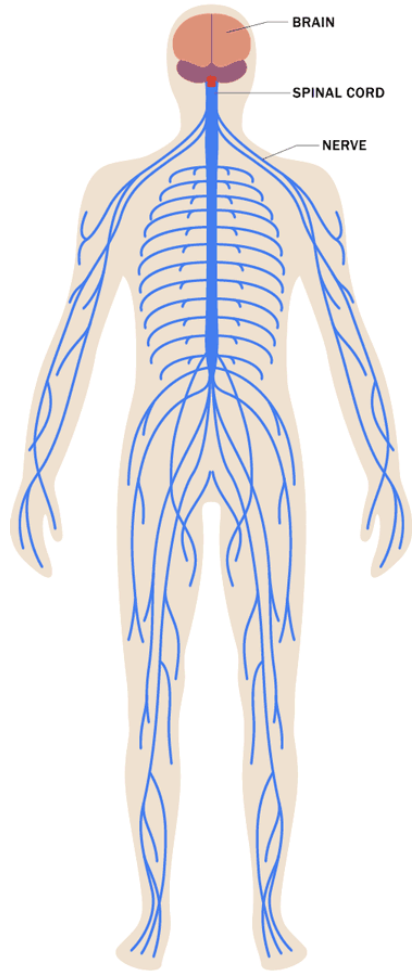

What is event streaming?

Technically speaking, event streaming is the practice of capturing data.

- *Sensors*
- *Mobile devices*
- *Cloud services*
- *Applications*

What is Apache Kafka?

What is Apache Kafka?

- An open-source distributed event streaming platform.
- Has three key capabilities:

1. To **publish** (write) and **subscribe to** (read) streams of events.
1. To **store** streams of events durably and reliably for as long as you want.
1. To **process** streams of events as they occur or retrospectively.

Why Kafka?

Why Kafka?

1. Fast
2. Scalable
3. Reliable
4. Higher throughput
5. Replication characteristics

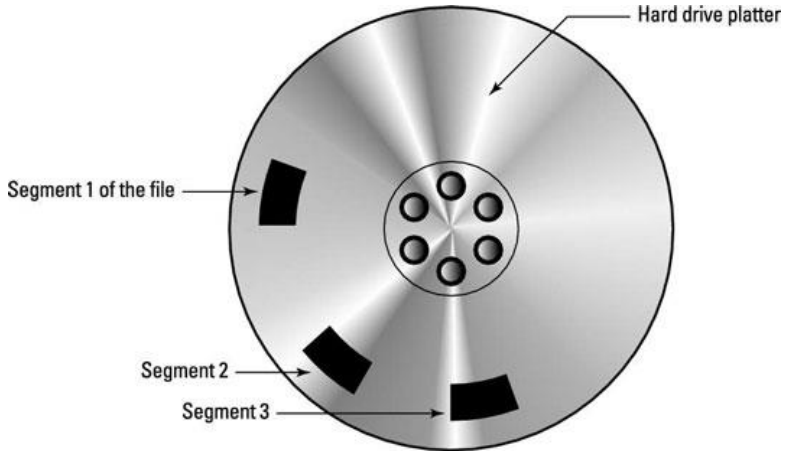

The most important reason Kafka is popular is Kafka's exceptional performance.

Delivery Guarantees & Throughput

Delivery Guarantees

- ***At most once*** which can lead to messages being lost, but they cannot be redelivered or duplicated
- ***At least once*** ensures messages are never lost but may be duplicated
- ***Exactly once*** guarantees each message processing (not delivery) happens once and only once

Throughput

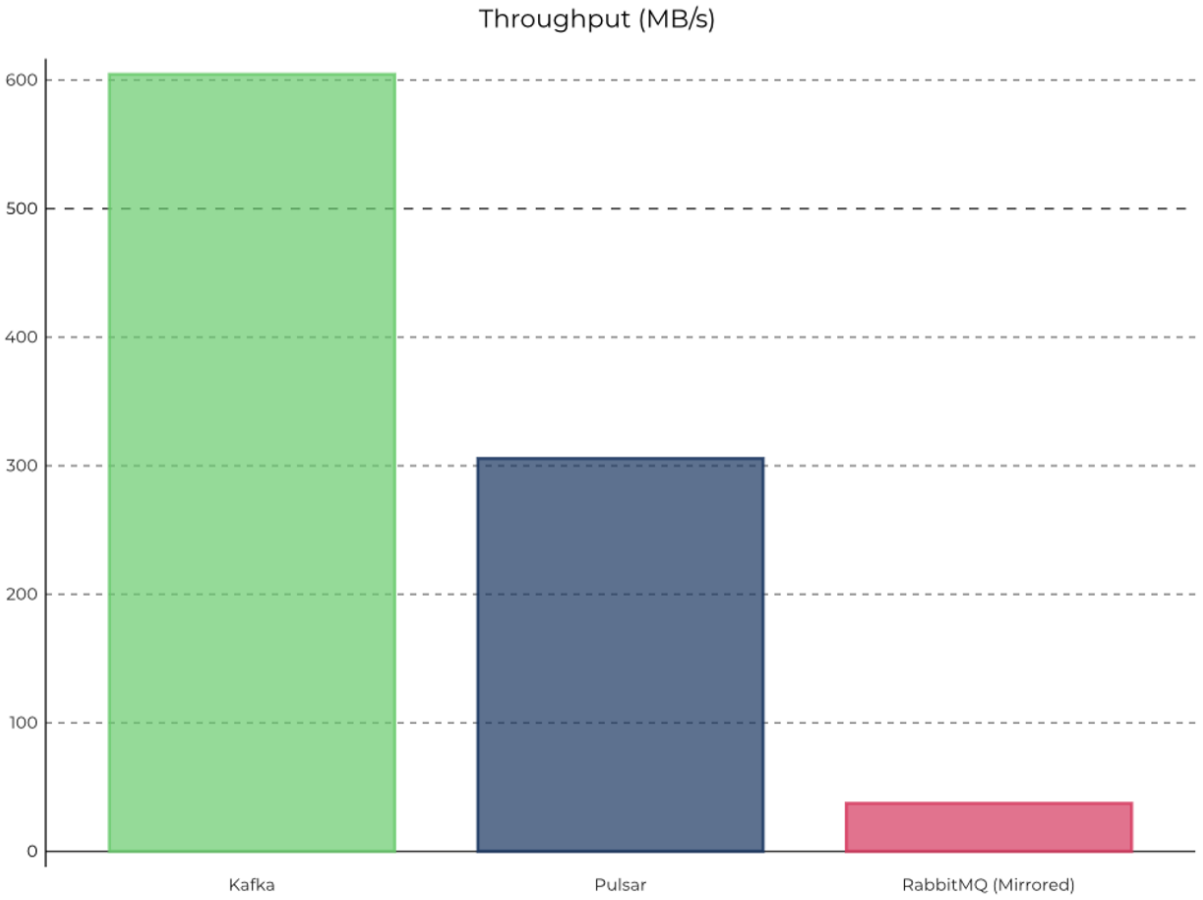

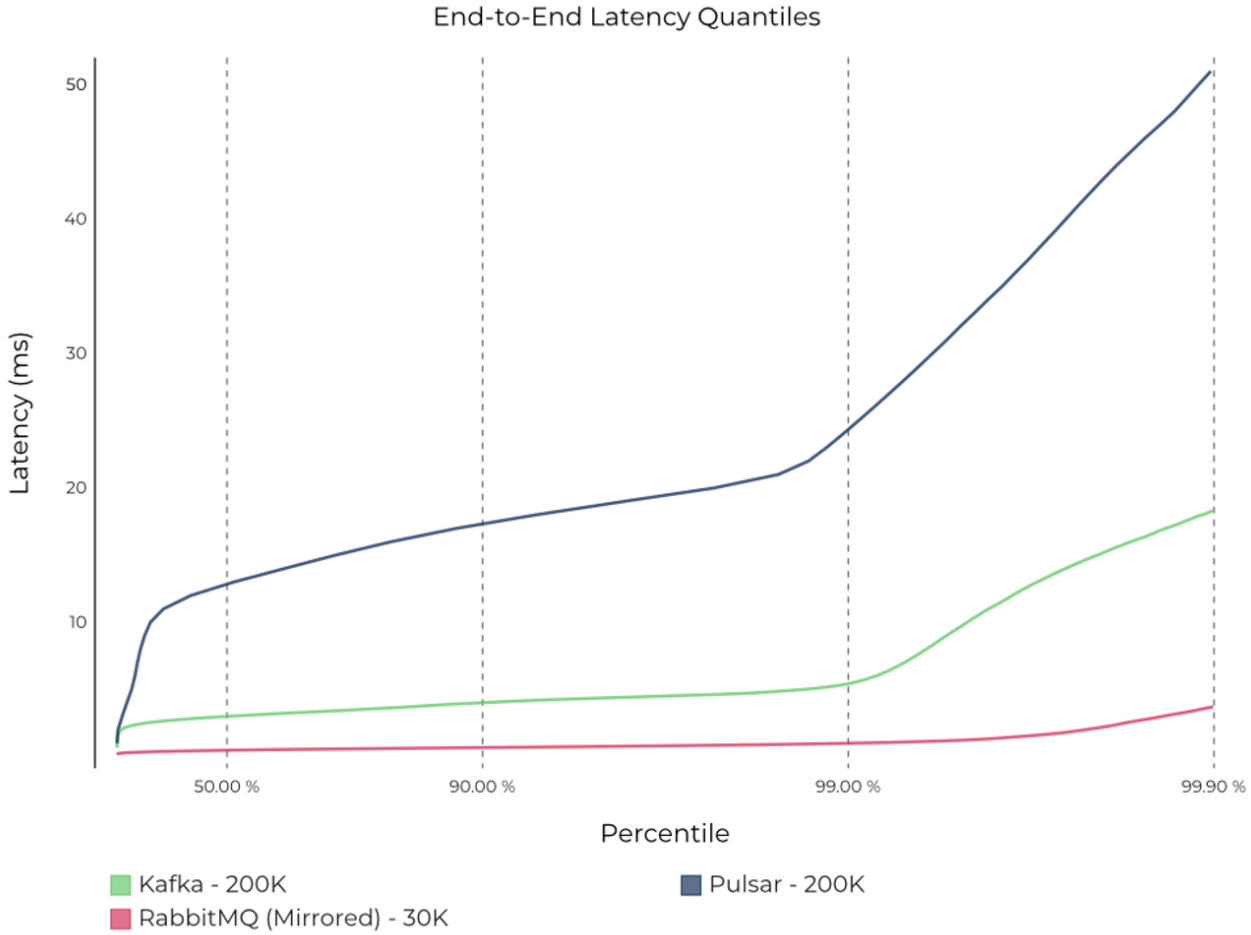

Throughput

	Kafka	Pulsar	RabbitMQ (Mirrored)
Peak Throughput (MB/s)	605 MB/s	305 MB/s	38 MB/s
p99 Latency (ms)	5 ms (200 MB/s load)	25 ms (200 MB/s load)	1 ms* (reduced 30 MB/s load)

**RabbitMQ latencies degrade significantly at throughputs higher than the 30 MB/s. Furthermore, the impact of mirroring is significant at higher throughput and better latencies can be achieved by using just classic queues without mirroring.*

Architecture and Components

Cluster Architecture

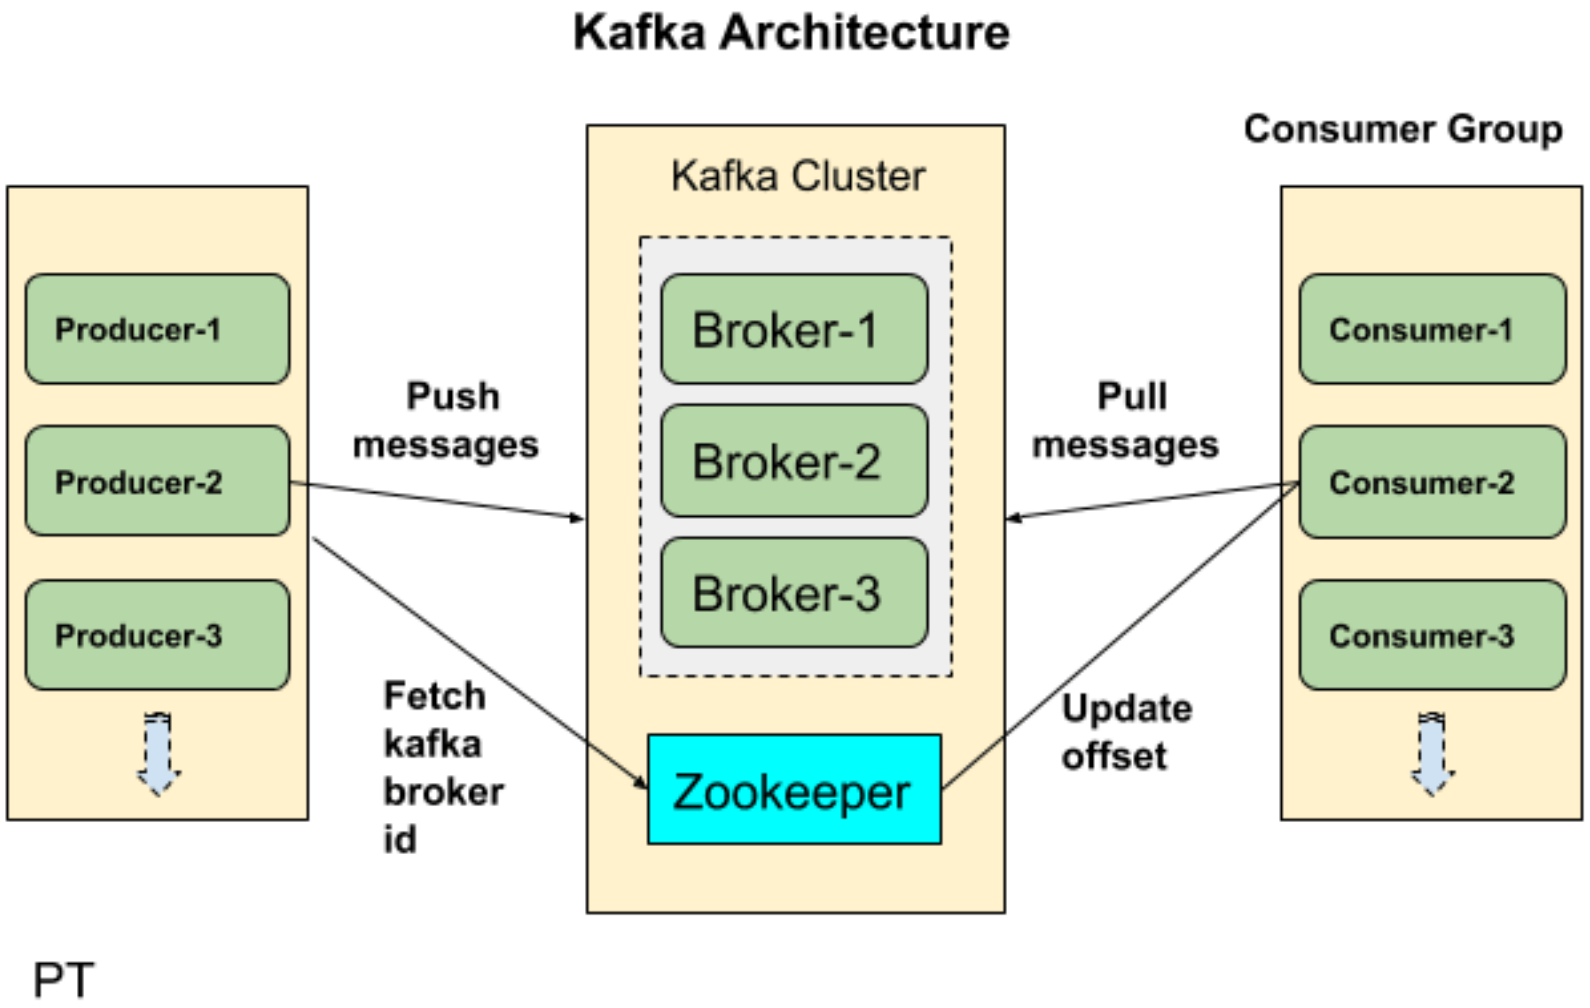

Broker

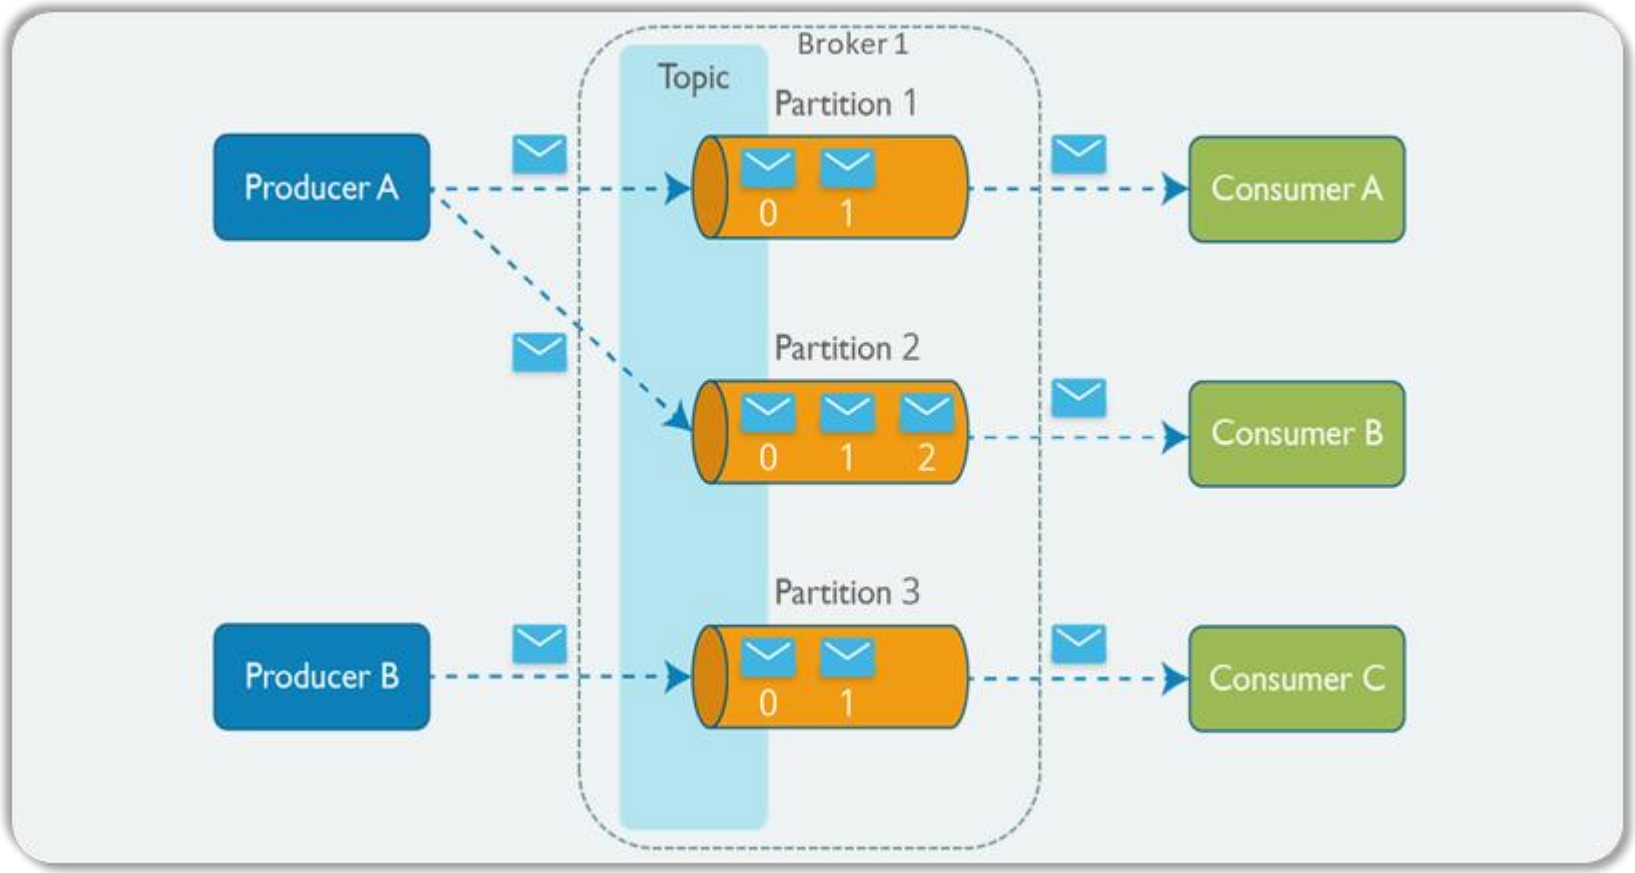

Topic Partition

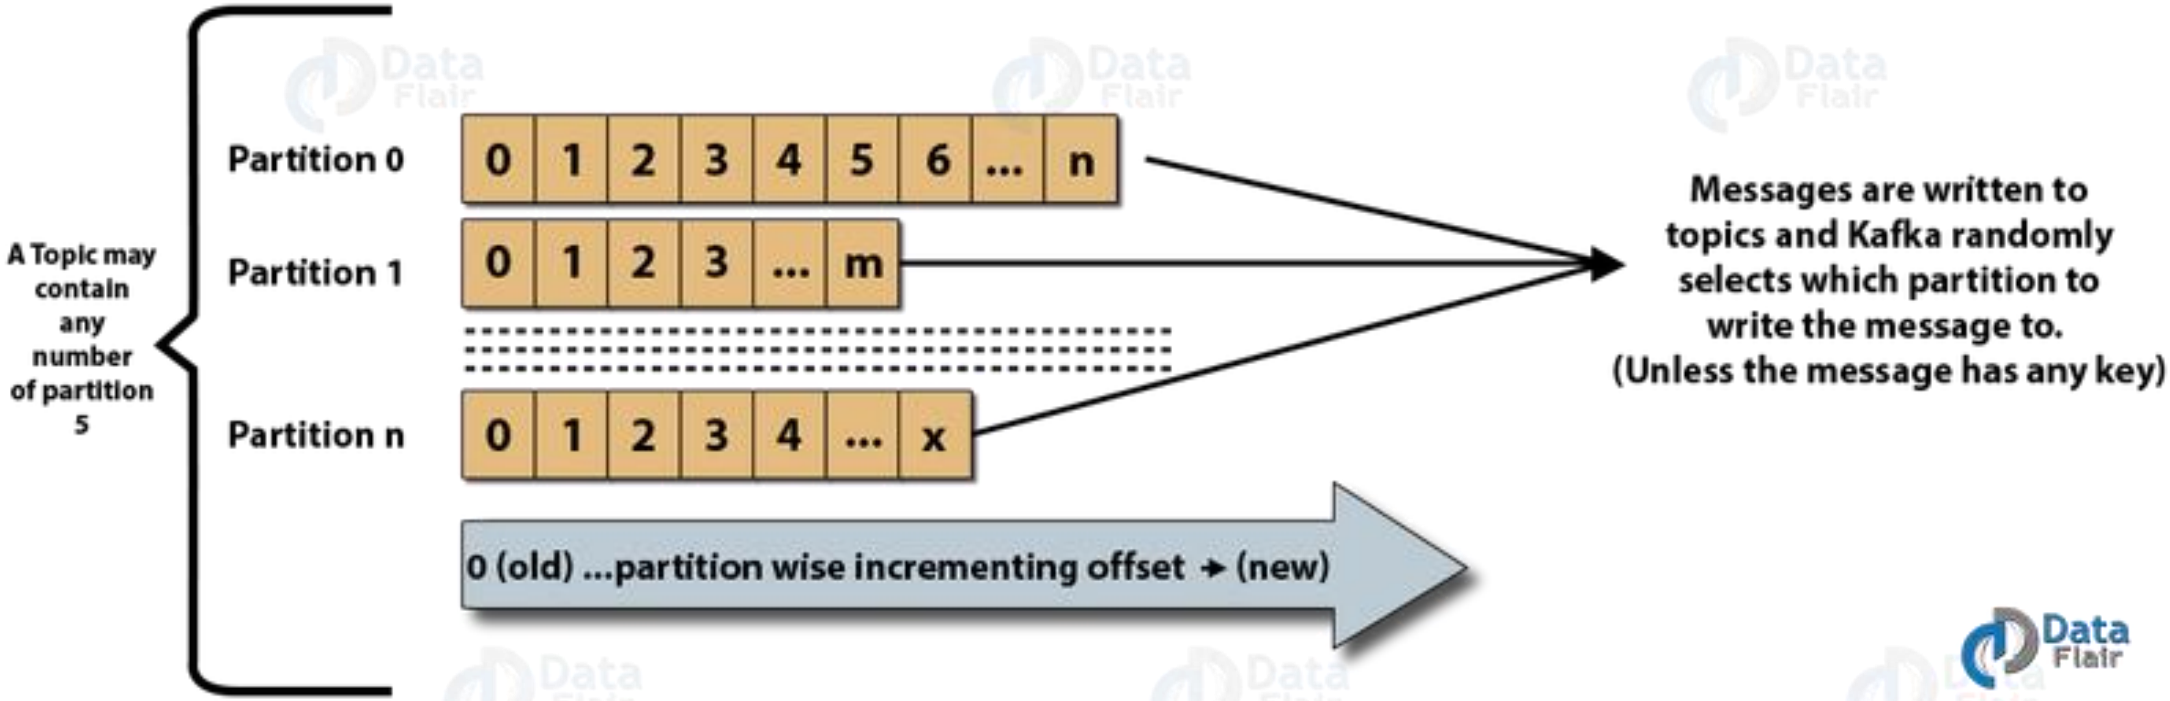

Topic replication factor

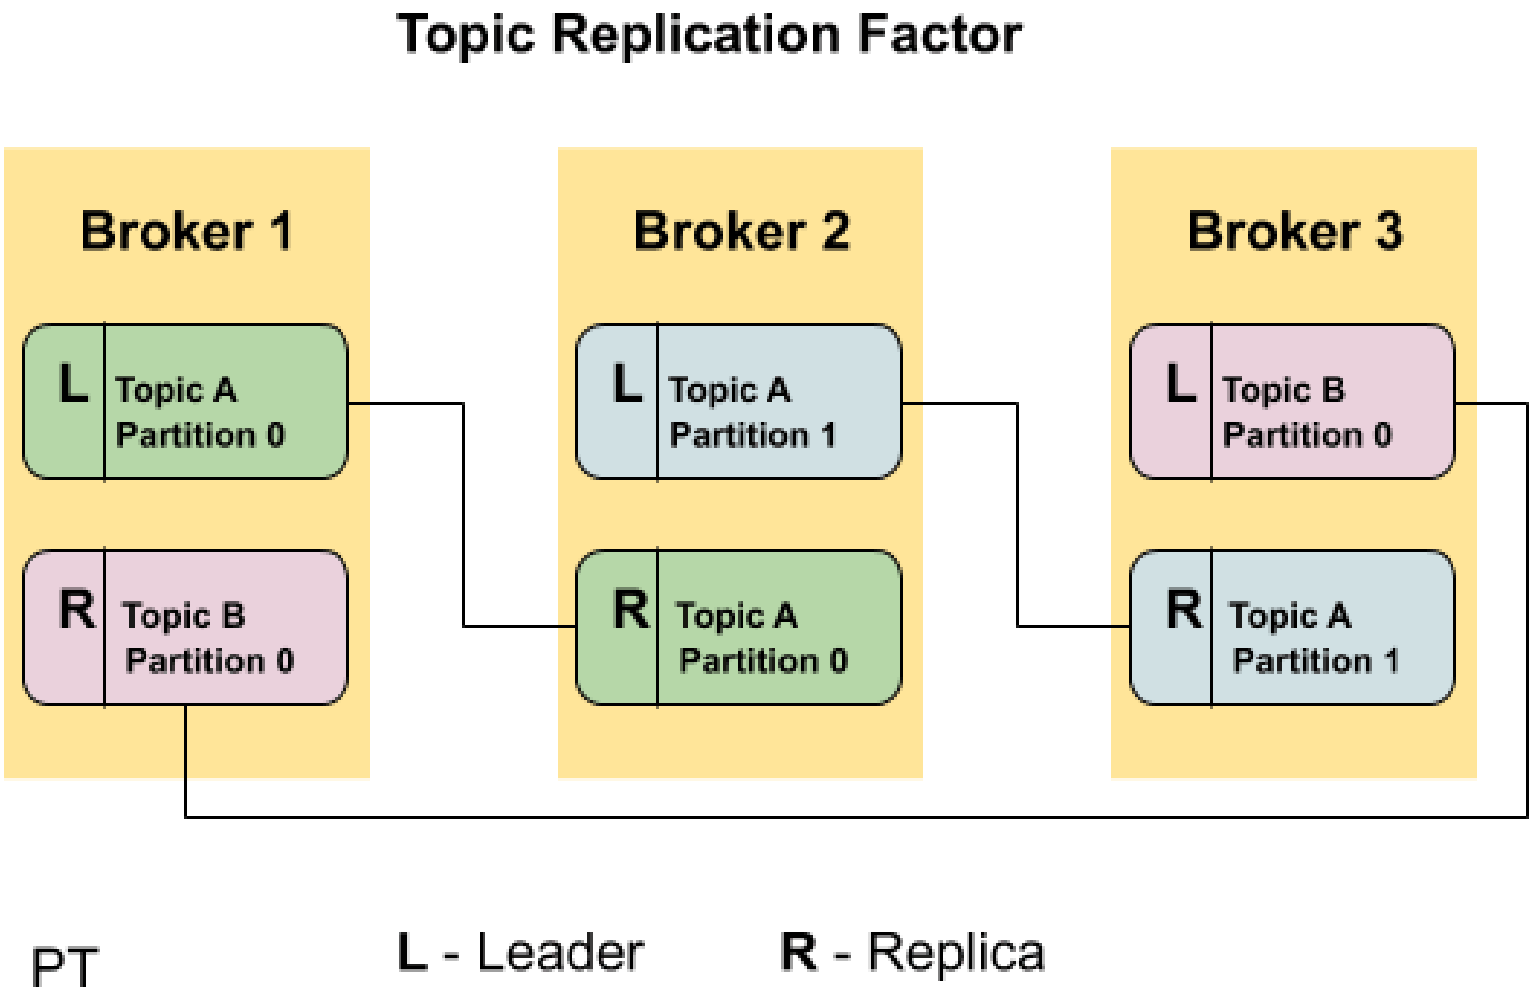

Topic configuration:

replication-factor=2, **partition**=10

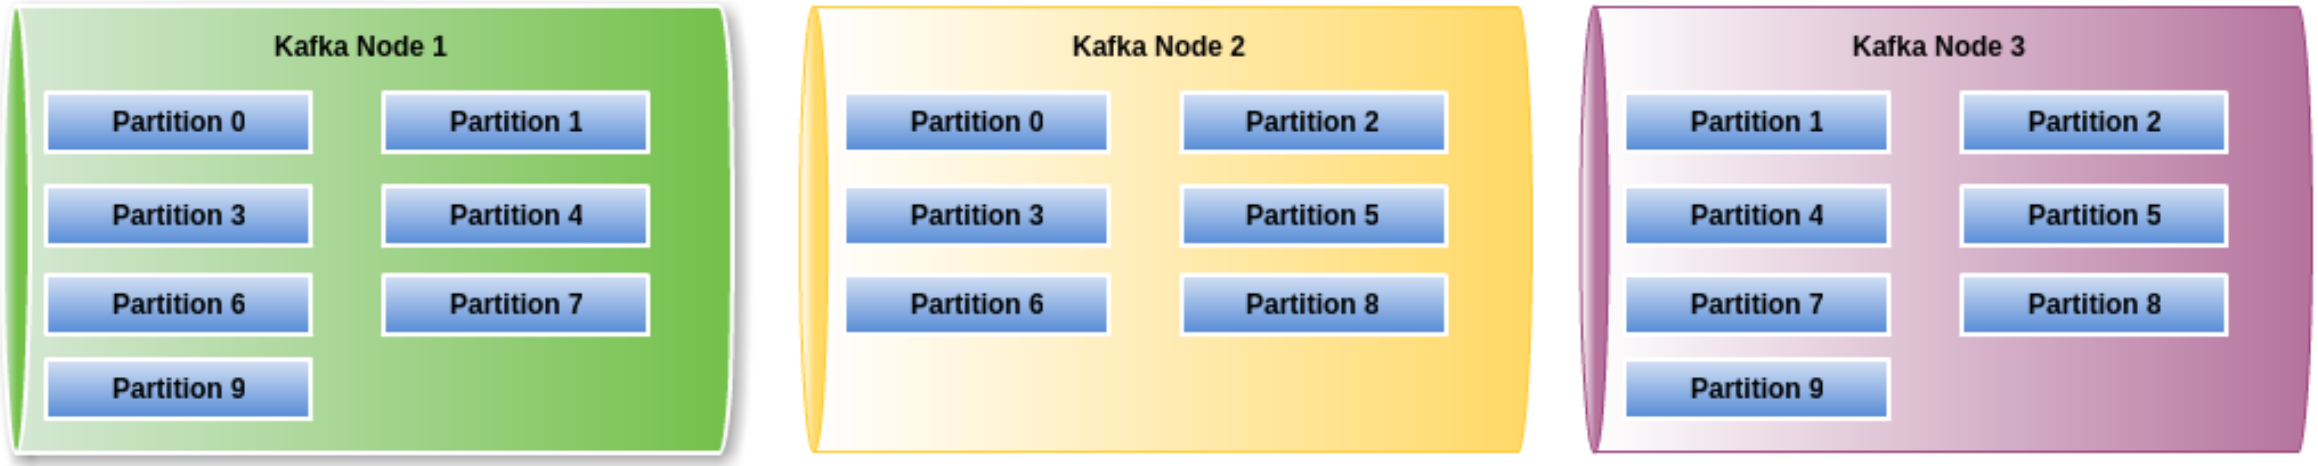

Topic partition configuration:

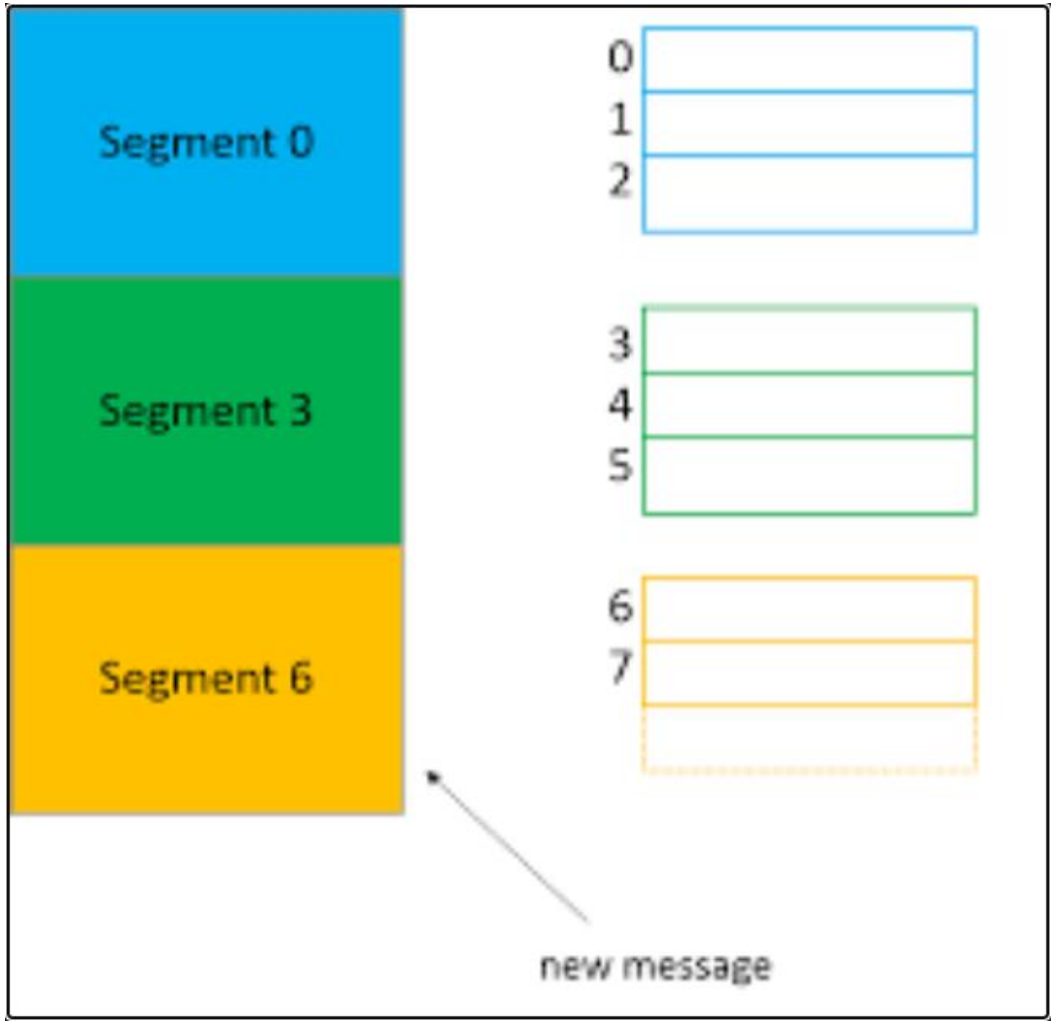

Consumer Group

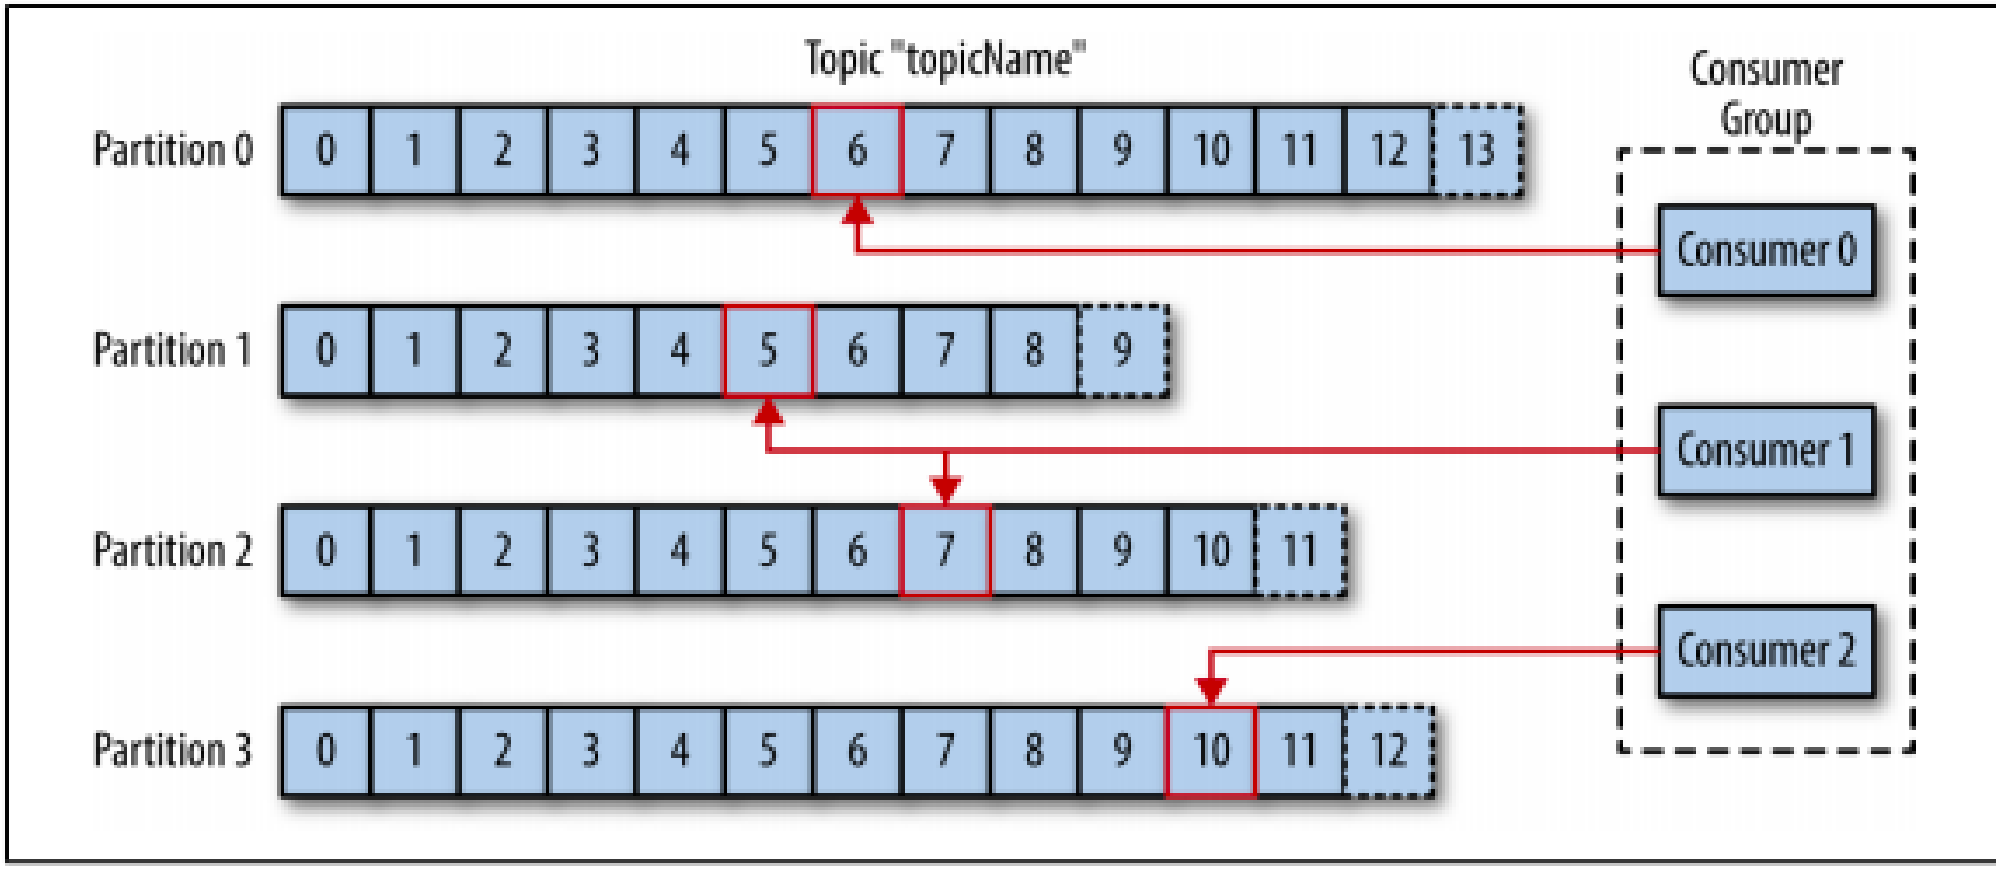

Consumer Group

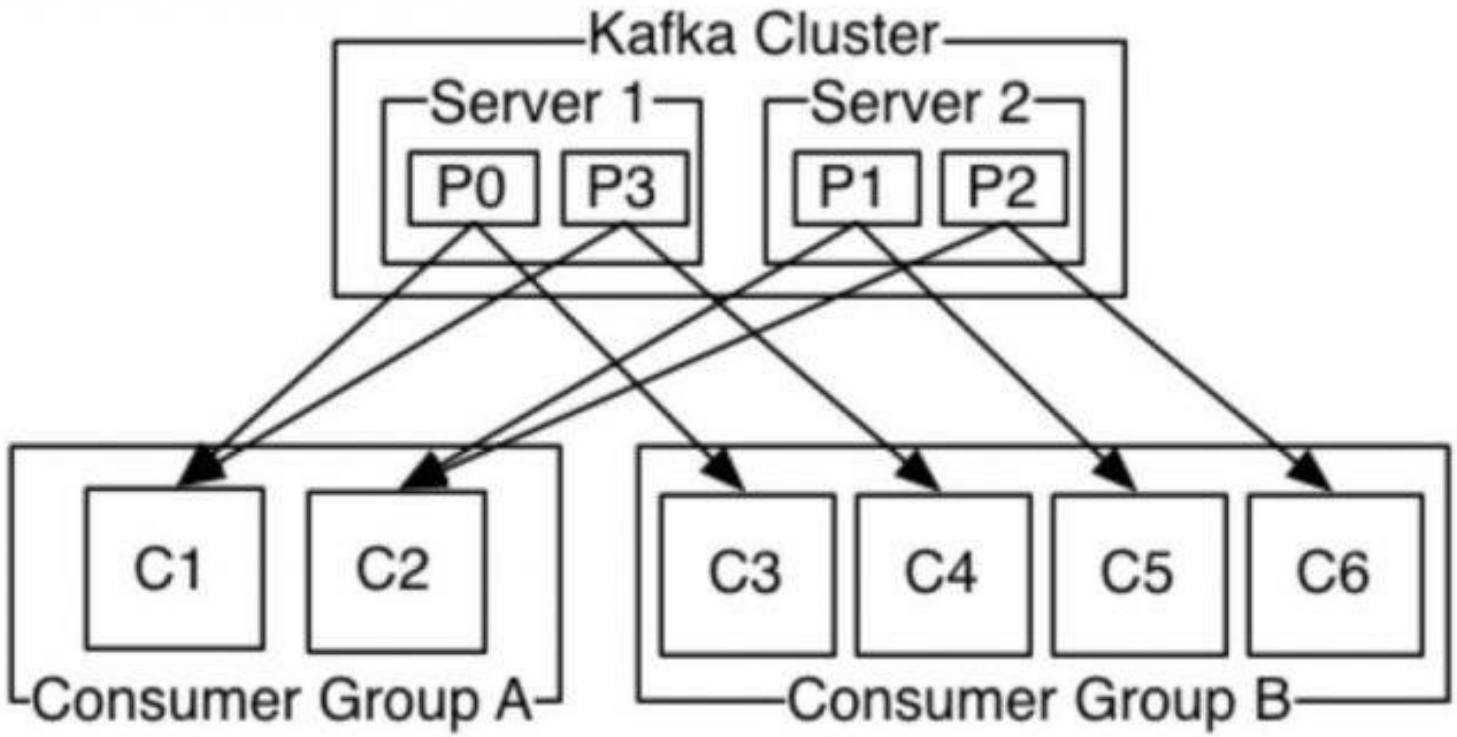

Kafka Connect

Kafka Connect

- A tool for scalably and reliably streaming data between Apache Kafka and other systems.
- Move large collections of data into and out of Kafka

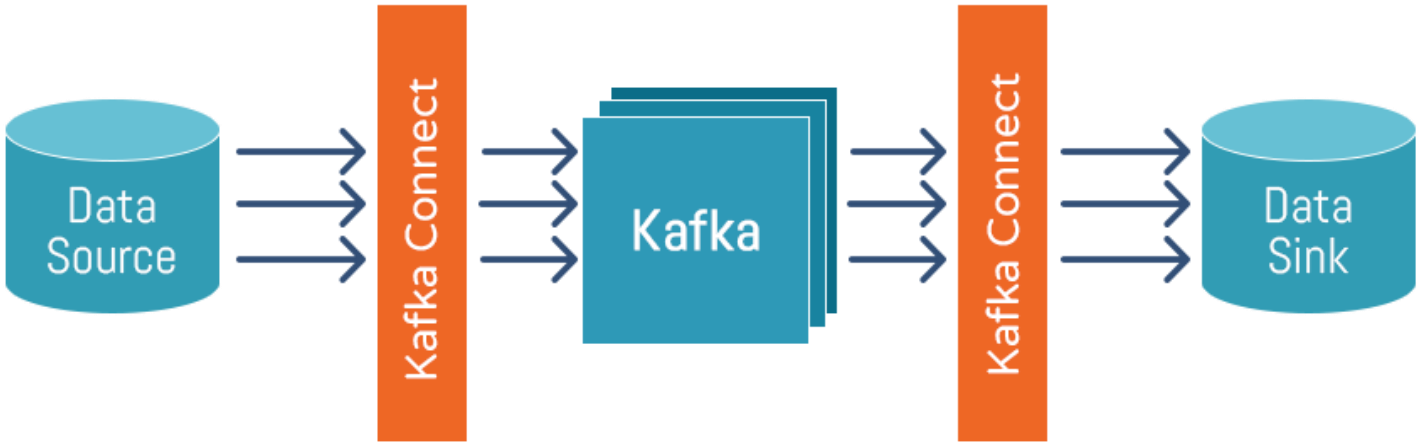

Kafka Stream

Kafka Stream

- A client library for building applications and microservices, where the input and output data are stored in Kafka clusters.

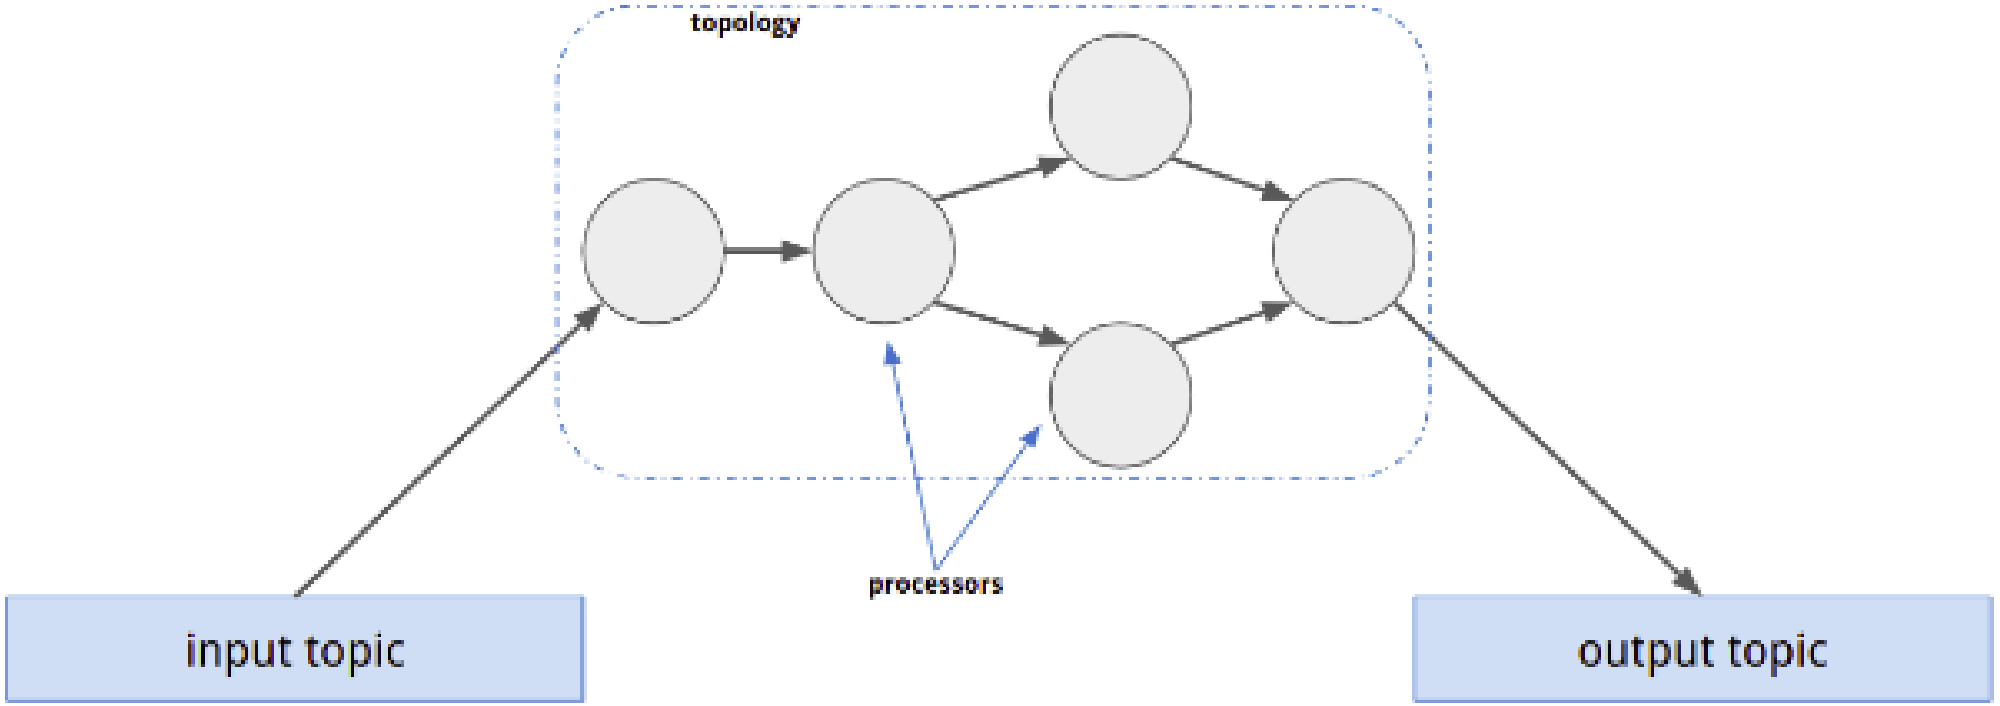

Stream DSL

- The Kafka Streams DSL (Domain Specific Language) is built on top of the Streams Processor API.
- **KStream**: *is an abstraction of a **record stream.***
- **KTable**: *is an abstraction of a **changelog stream.***
- **GlobalKTable**: *is an abstraction of a **changelog stream**.*

Security

Security

- Kafka supports the following SASL mechanisms:
  SASL/GSSAPI (Kerberos)
  SASL/PLAIN
  SASL/SCRAM-SHA-256 and SASL/SCRAM-SHA-512
  SASL/OAUTHBEARER

- Authentication of connections from brokers to ZooKeeper

- Encryption of data transferred between brokers and clients, between brokers.

- Authorization of read / write operations by clients

Kafka Pros-Cons

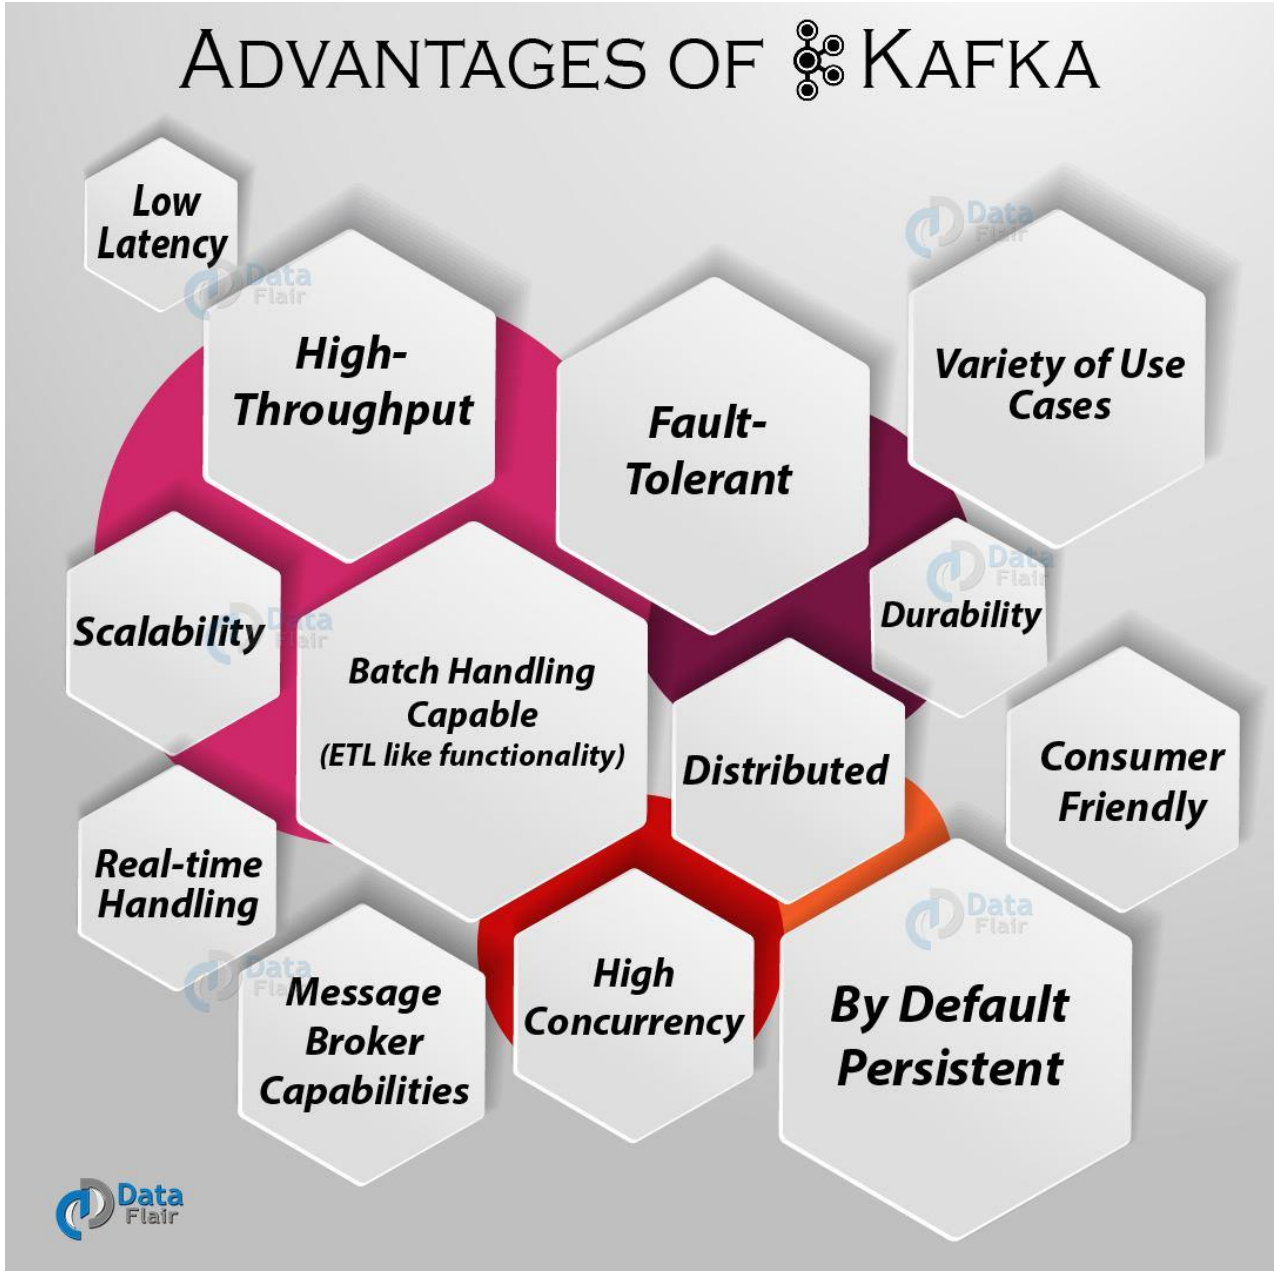

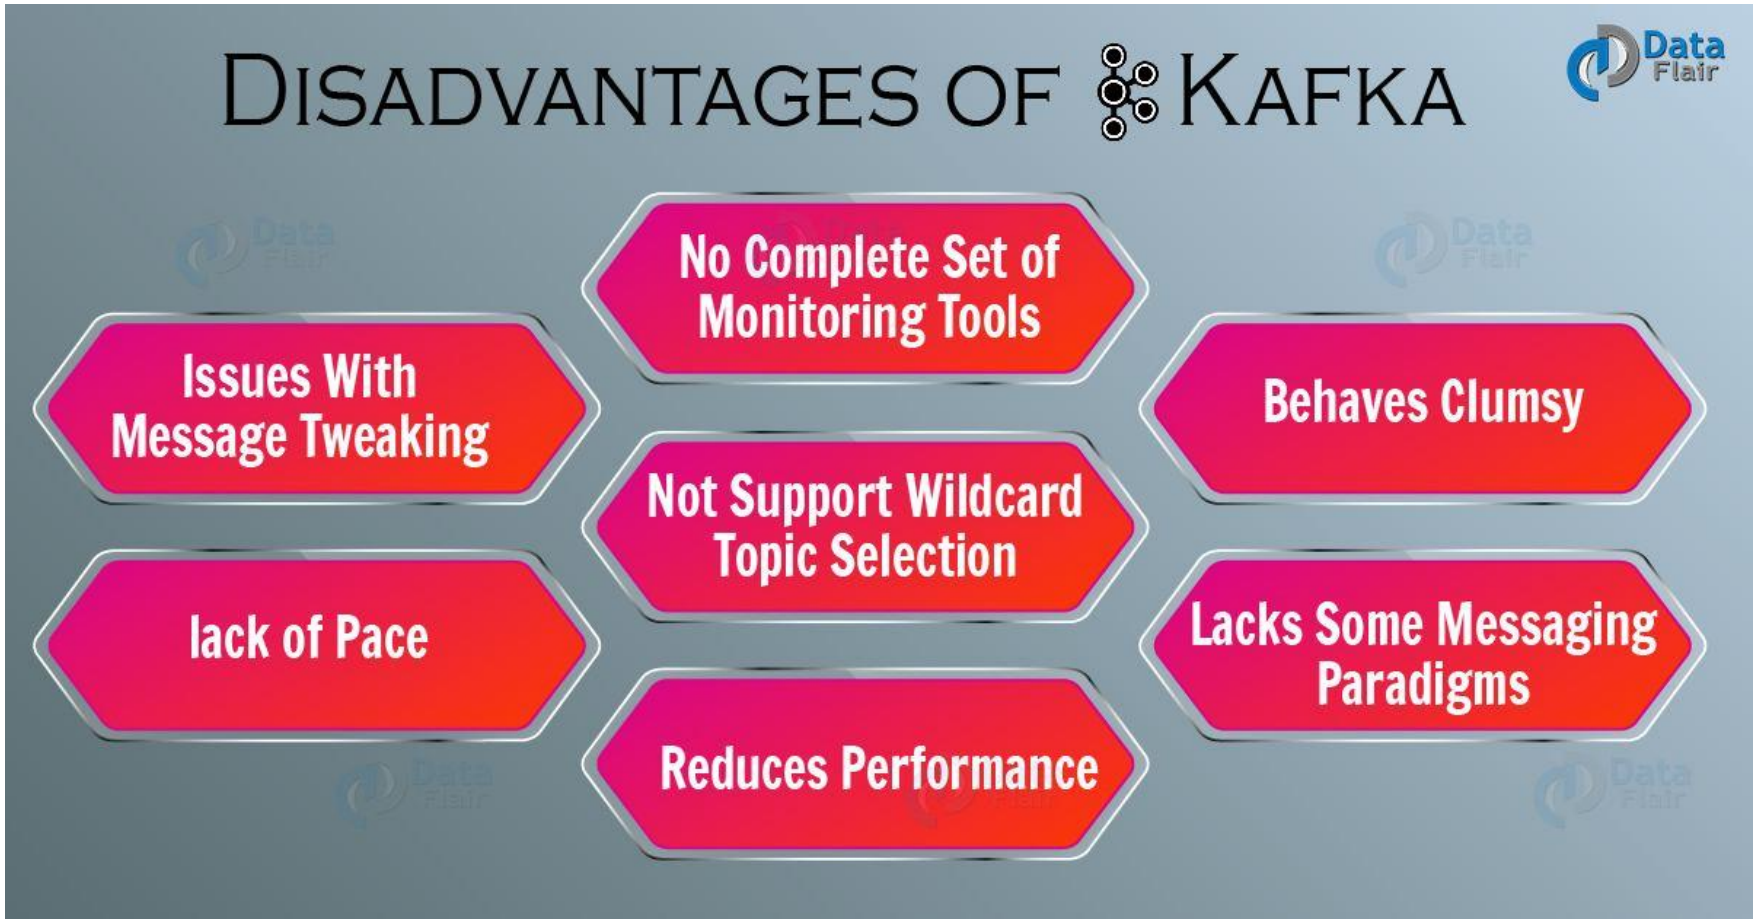

Demo

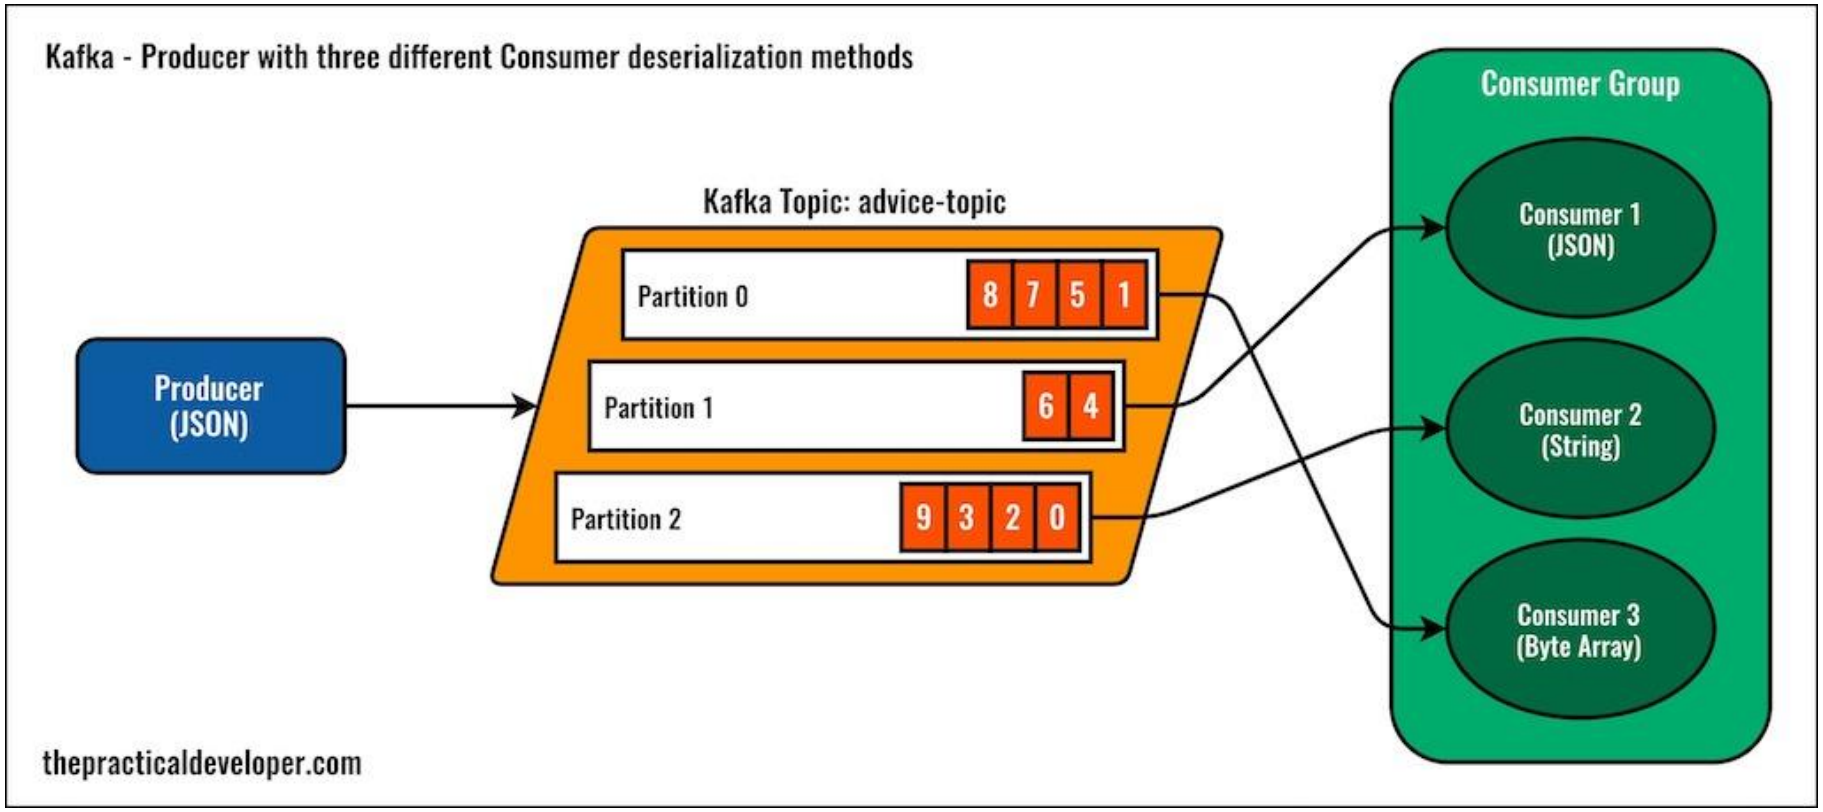

References

- https://kafka.apache.org/documentation
- https://kafka.apache.org/20/documentation/streams/developer-guide/dsl-api.html
- https://docs.confluent.io/platform/current/ksqldb/index.html
- https://dzone.com/articles/what-is-kafka#:~:text=Why%20Kafka%3F,due%20to%20volume%20and%20responsiveness

Thanks